

# DATOMIC VS. CRUX AND WHY IT MATTERS

@tiagoluchini

<https://luchini.nyc>



# TIAGO LUCHINI

@tiagoluchini  
<https://luchini.nyc>







ANO II - Nº 13  
OUTUBRO 1982  
Cr\$ 350,00  
ISSN 0101-3401

# Micro Sistemas

A PRIMEIRA REVISTA BRASILEIRA DE MICROCOMPUTADORES

UM ANO  
DE MICRO SISTEMAS



VisiCalc  
Tecnologia  
aberta no DEL

SORTs  
comparados

Conheça o  
6502  
Impressoras

ANO II - Nº 13  
OUTUBRO 1982  
Cr\$ 350,00

ISSN 0101 - 3401

# Micro Sistemas

A PRIMEIRA REVISTA BRASILEIRA DE MICROCOMPUTADORES

UM ANO  
DE MICRO SISTEMAS



VisiCalc  
Tecnologia  
charta no DFI

SORTs  
comparados

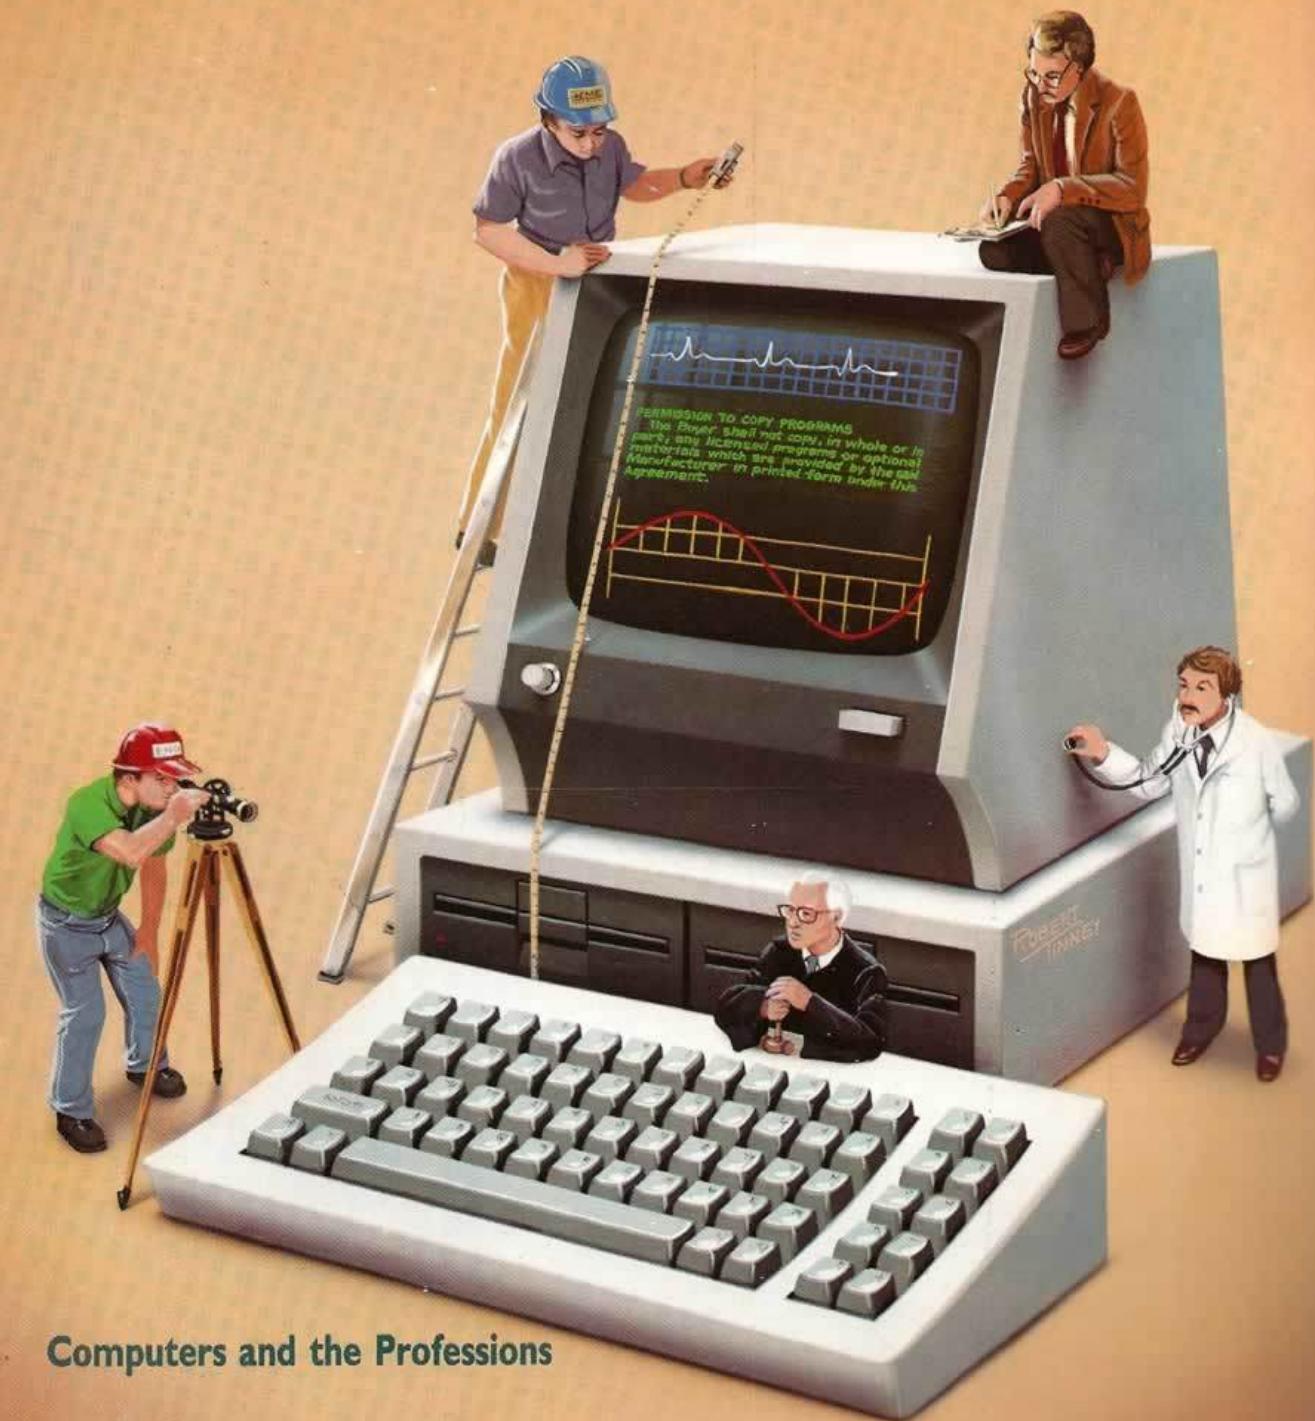
Conheça o  
6502  
Impressoras

12490

MAY 1984 Vol. 9, No. 5  
\$3.50 in U.S.A.  
\$4.25 in Canada/£2.10 in U.K.  
A McGraw-Hill Publication  
0360-5280

# BYTE

the small systems journal



Computers and the Professions

```

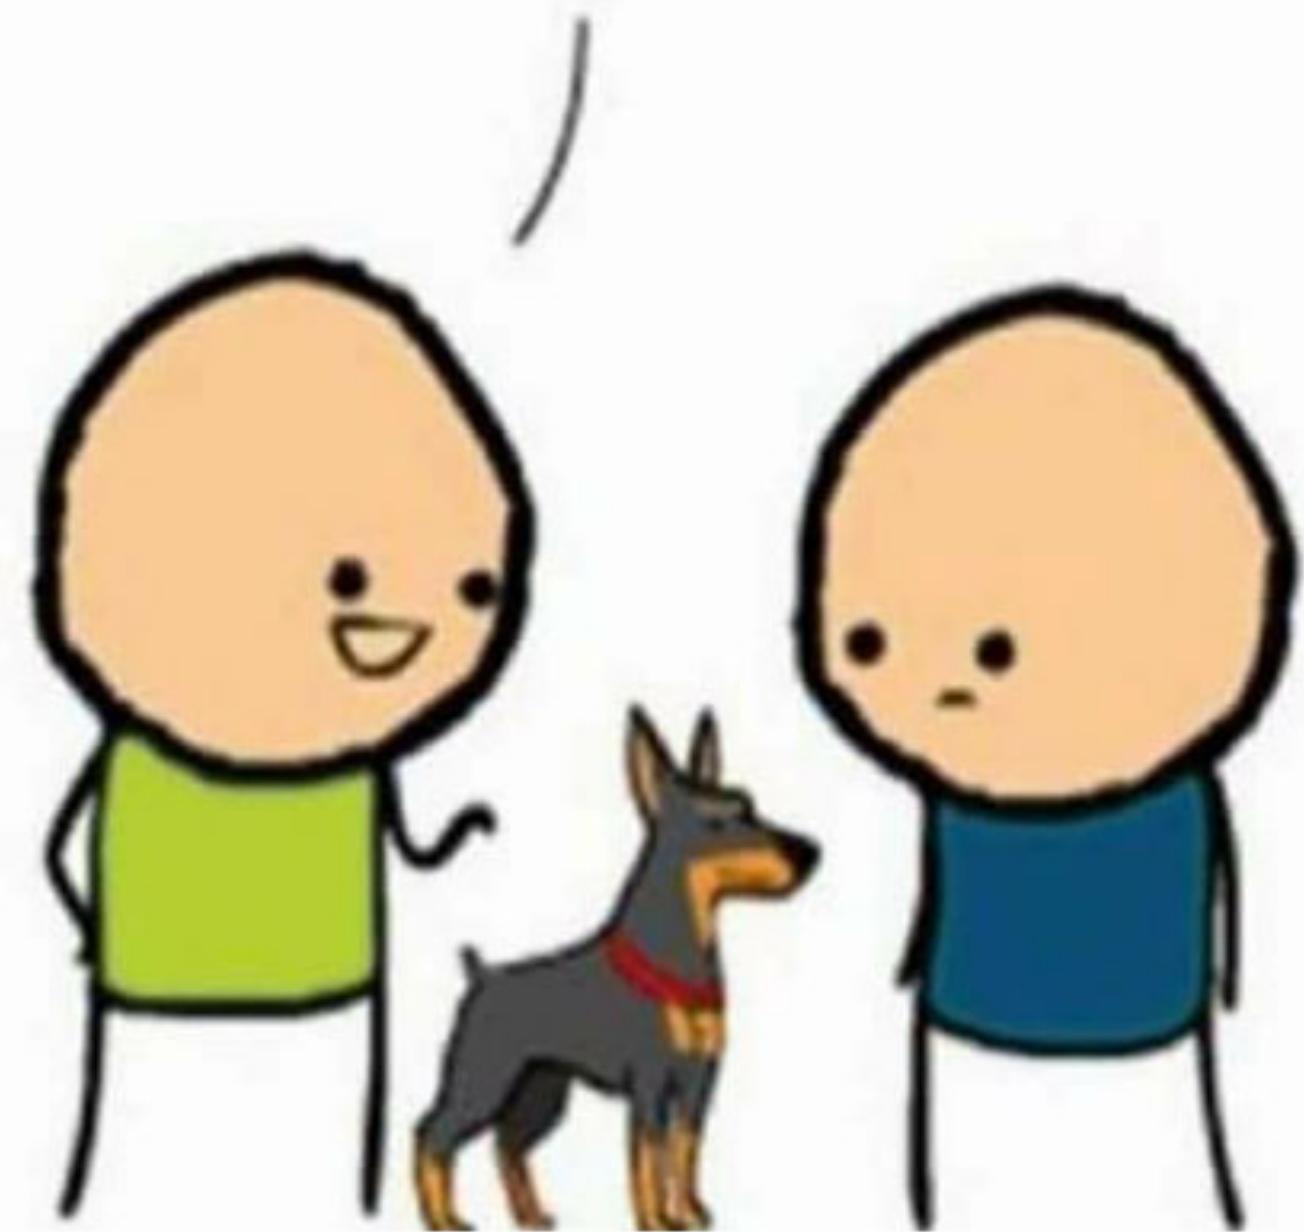
670 INPUT *      ESTA TUDO CORRETO (S/N) *;S$*
671 IF S$!="N" THEN GOTO 628
672 IF SNOME<>"* THEN MAT#(I,1)=SNOME#
673 IF NOME<>"* THEN MAT#(I,2)=NOME#
674 IF ENDE<>"* THEN MAT#(I,3)=ENDE#
675 IF TEL<>"* THEN MAT#(I,4)=TEL#
676 MD#=1
677 NEXT I
678 IF MD#>1 GOTO 683
679 PRINT *
680 PRINT *      SOBRENOME INEXISTENTE*
681 FOR K=1 TO 1000
682 NEXT K
683 RETURN
700 REM -----< OPCAO PESQUISA >
701 PI=0
702 EI=0
703 HOME
704 PRINT *      PESQUISA*
705 VTAB 6
706 INPUT *      ENTRE COM O SOBRENOME *;SNOME#
707 SNOME#=SNOME#+SPACE$(N11-LEN(SNAME#))
708 FOR I=2 TO NMAX#1
709 IF SNOME$<>MAT#(I,1) OR MAT#(I,5)<>"1" THEN GOTO 732
710 EI=1
711 HOME
712 PRINT *      PESQUISA*
713 VTAB 7
714 PRINT
715 PRINT "SOBRENOME ";MAT#(I,1)
716 PRINT
717 PRINT "NOME   ";MAT#(I,2)
718 PRINT
719 PRINT "ENDERECO  ";MAT#(I,3)
720 PRINT
721 PRINT "TELEFONE  ";MAT#(I,4)
722 PRINT
723 VTAB 21
724 INPUT *      ESTE E' O REGISTRO PEDIDO (S/N) *;SN#
725 IF SN#!="N" OR SN#!="S" GOTO 729
726 PRINT
727 INPUT *      ENTRE S OU N *;SN#
728 GOTO 725
729 IF SN#!="N" GOTO 732
730 PI=1
731 GOTO 733
732 NEXT I
733 IF EI=1 GOTO 739
734 PRINT
735 PRINT *      SOBRENOME INEXISTENTE*
736 FOR K=1 TO 1000
737 NEXT K
738 RETURN
739 IF PI=1 GOTO 745
740 IF PI<>0 OR EI<>1 GOTO 745
741 PRINT
742 PRINT *      NAO EXITE MAIS NENHUM ";SNOME#
743 FOR K=1 TO 1000
744 NEXT K
745 RETURN
800 REM -----< OPCAO LISTAGEM >
801 REM
802 REM -----< TELA DE OPCOES PARA CLASSIFICACAO >
803 REM
804 HOME
805 PRINT *      LISTAGEM*
806 VTAB 9
807 PRINT *      ESCOLHA A CHAVE DE CLASSIFICACAO*
808 PRINT
809 PRINT *      SOBRENOME (1)*
810 PRINT *      NOME    (2)*
811 PRINT *      ENDERECO (3)*
812 PRINT *      TELEFONE (4)*
813 PRINT
814 INPUT *      ENTRE COM A OPCAO *;OP#
815 IF OP#>1 AND OP#<=4 GOTO 821
816 PRINT
817 PRINT *      CHAVE INEXISTENTE*
818 FOR K=1 TO 1000
819 NEXT K
820 GOTO 804
821 IF OP#>1 GOTO 824
822 NCHAR#=N1#
823 INIC#=1
824 IF OP#>2 GOTO 827
825 NCHAR#=N2#
826 INIC#=11
827 IF OP#>3 GOTO 830
828 NCHAR#=N3#
829 INIC#=41
830 IF OP#>4 GOTO 833
831 NCHAR#=N4#
832 INIC#=81
833 J=1
834 FOR I=2 TO NMAX#1
835 IF MAT#(I,5)<>"1" GOTO 845
836 A#=MAT#(I,1)+SPACE$(N11-LEN(MAT#(I,1)))
837 B#=MAT#(I,2)+SPACE$(N21-LEN(MAT#(I,2)))
838 C#=MAT#(I,3)+SPACE$(N31-LEN(MAT#(I,3)))
839 D#=MAT#(I,4)+SPACE$(N41-LEN(MAT#(I,4)))
840 E#=MAT#(I,5)
841 R#(J)=A#+B#+C#+D#+E#
842 J=J+1
843 NEXT I
844 NREG#=J-1
845 REM -----< OPCAO CLASSIFICACAO >
846 GOSUB 900
847 FOR I=2 TO NTOTAL#1
848 MAT#(I,1)=MID$(R#(I-1),1,N1#)
849 MAT#(I,2)=MID$(R#(I-1),N1#+1,N2#)
850 MAT#(I,3)=MID$(R#(I-1),N1#+N2#+1,N3#)
851 MAT#(I,4)=MID$(R#(I-1),N1#+N2#+N3#+1,N4#)
852 MAT#(I,5)=MID$(R#(I-1),N1#+N2#+N3#+N4#+1,1)
853 NEXT I
854 REM -----< LISTAGEM >
855 FOR I=2 TO NTOTAL#1
856 HOME
857 PRINT *      LISTAGEM*
858 VTAB 9
859 PRINT "SOBRENOME ";MAT#(I,1)
860 PRINT
861 PRINT "NOME   ";MAT#(I,2)
862 PRINT
863 PRINT "ENDERECO  ";MAT#(I,3)
864 PRINT
865 PRINT "TELEFONE  ";MAT#(I,4)
866 VTAB 23
867 INPUT *ENTRE RETURN PARA CONTINUAR *;AA#
868 NEXT I
869 RETURN
900 REM -----< SUBROTINA DE CLASSIFICACAO >
901 DEFINT A-J,L-Q
902 DEFINT S-Z
903 M=9
904 K#(0)=STRING$(NCHAR,0)
905 K#(NREG+1)=STRING$(NCHAR,127)
906 REM
907 FOR I=1 TO NREG
908 K#(I)=MID$(P#(I),INIC,NCHAR)
909 NEXT I
910 IF NREG<=0 THEN GOTO 961
911 IP#1
912 P#(IP,1)=0
913 P#(IP,2)=0
914 L=1
915 S=NREG
916 I=L
917 J=S+
918 KEY#=K#(L)
919 I=I+
920 IF K#(I)=KEY# THEN GOTO 919
921 J=J-1

```

DOES HE BITE?



NO, BUT HE CAN HURT  
YOU IN OTHER WAYS



# **CRUX & DATOMIC REASONS WHY YOU SHOULD CARE**

# **REASONS WHY YOU SHOULD CARE**

1. Immutable Database
2. Query Like a Ninja
3. Unbundle Your Database

# REASONS WHY YOU SHOULD CARE

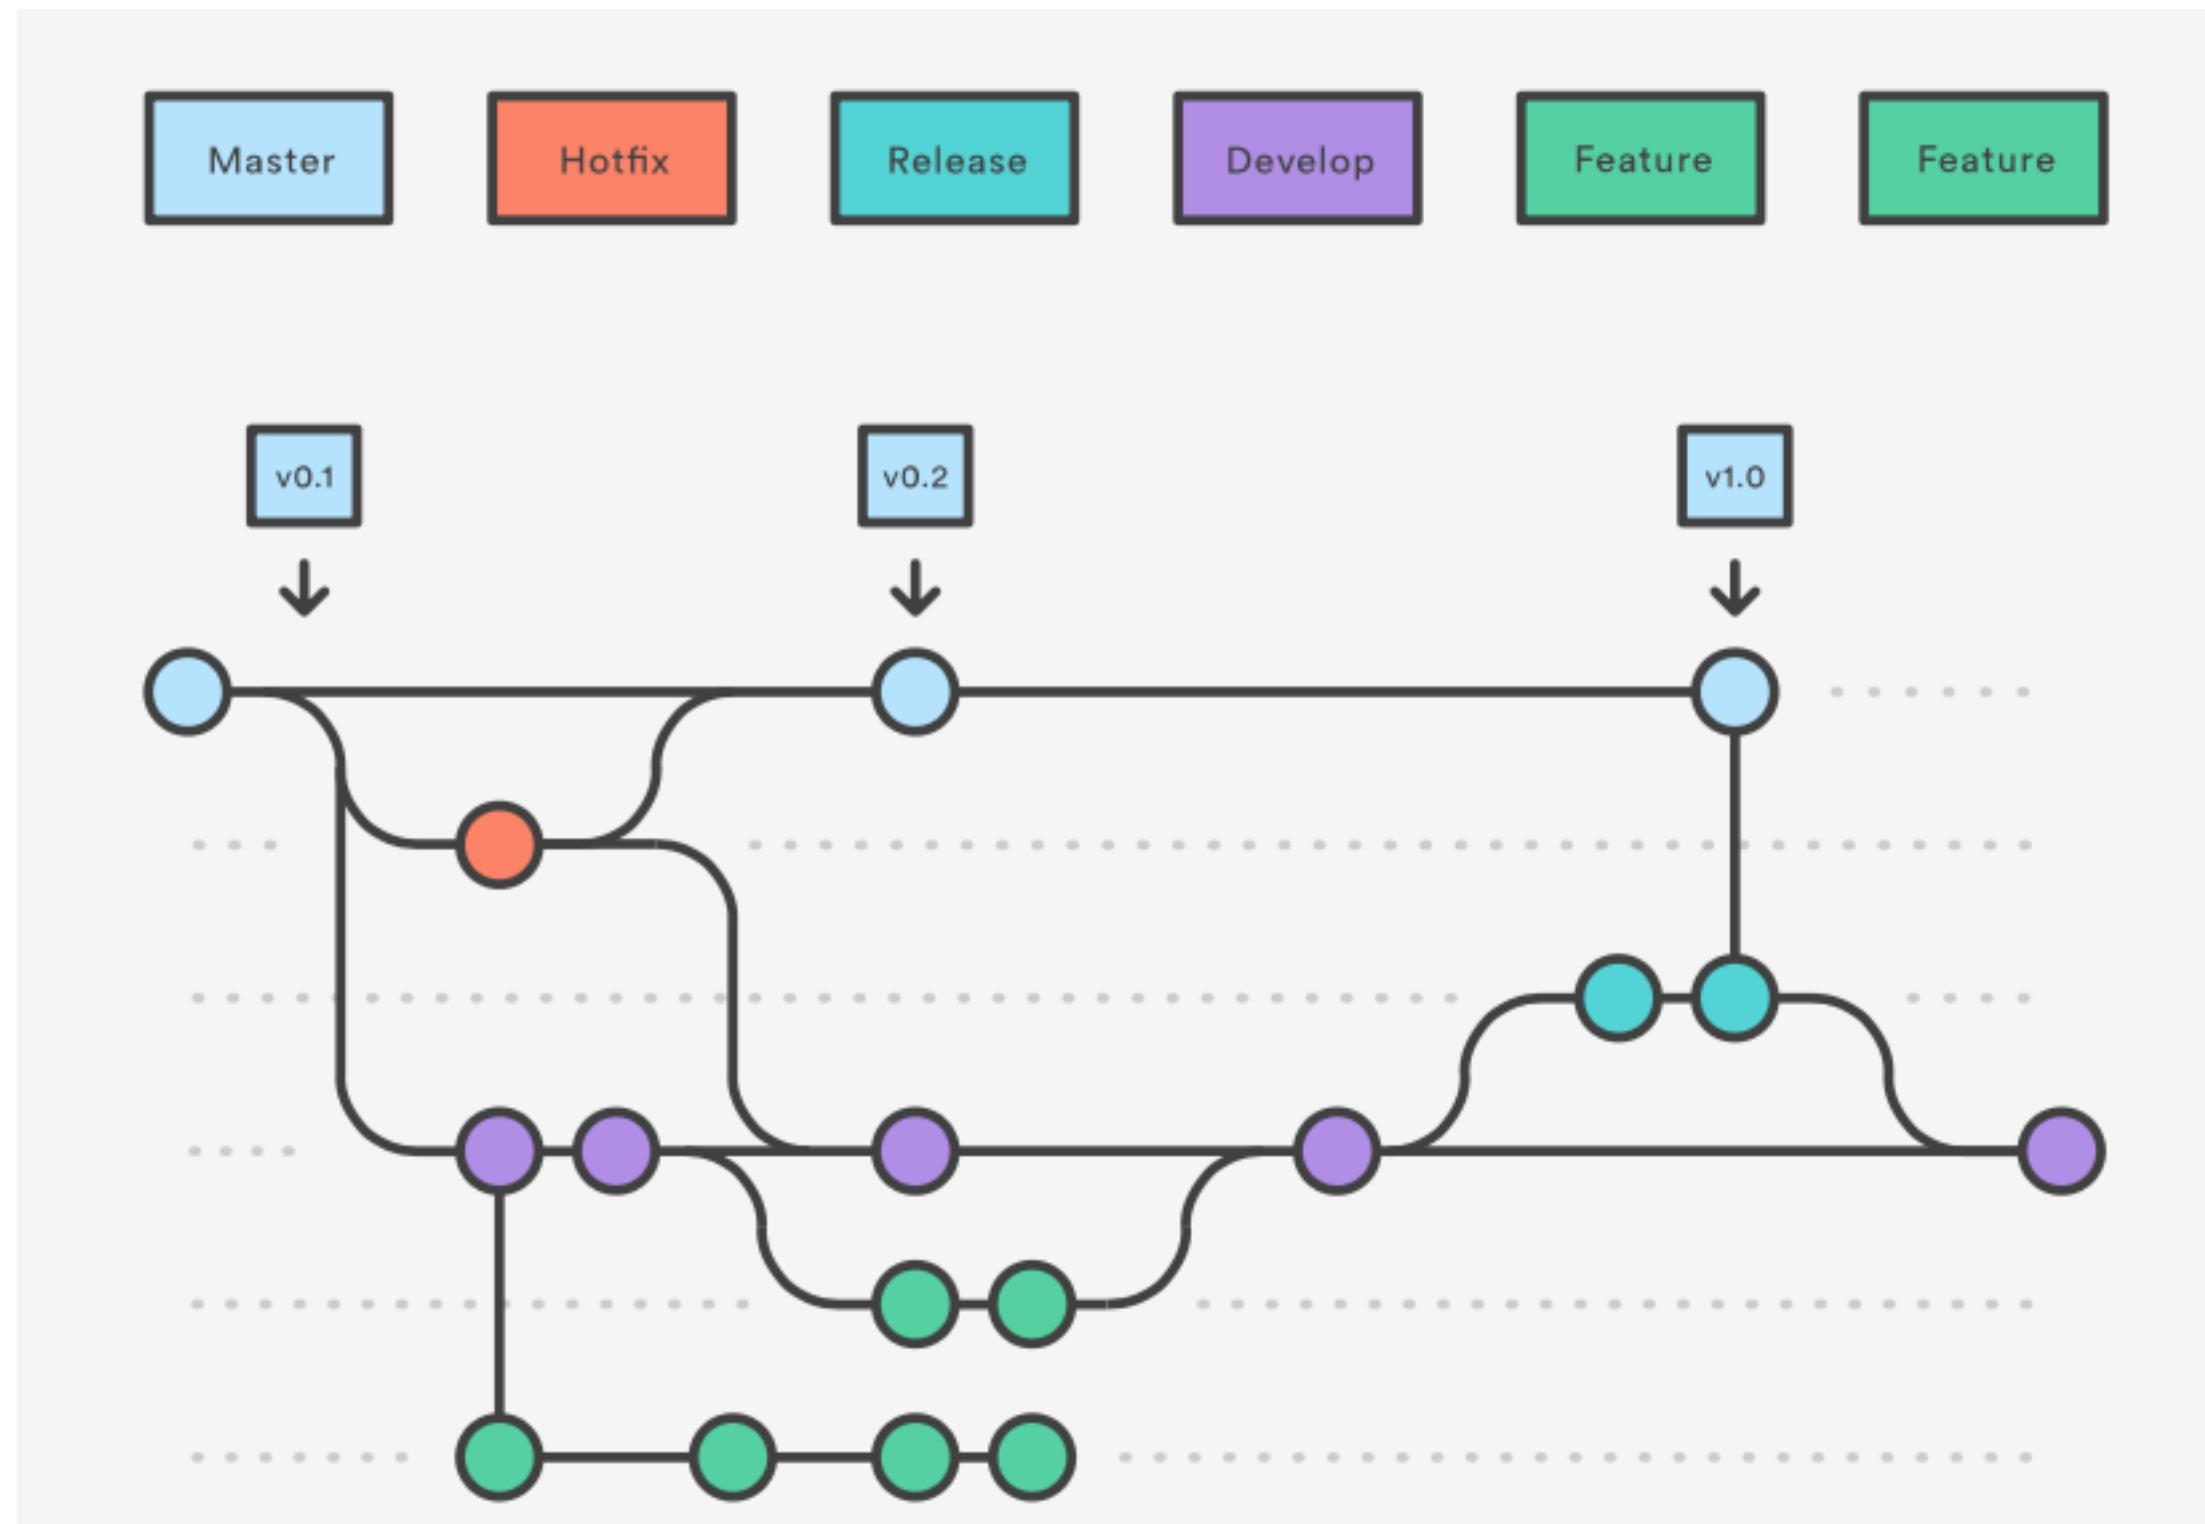
1. Immutable Database
2. Query Like a Ninja
3. Unbundle Your Database

# IMMUTABILITY EVERYWHERE

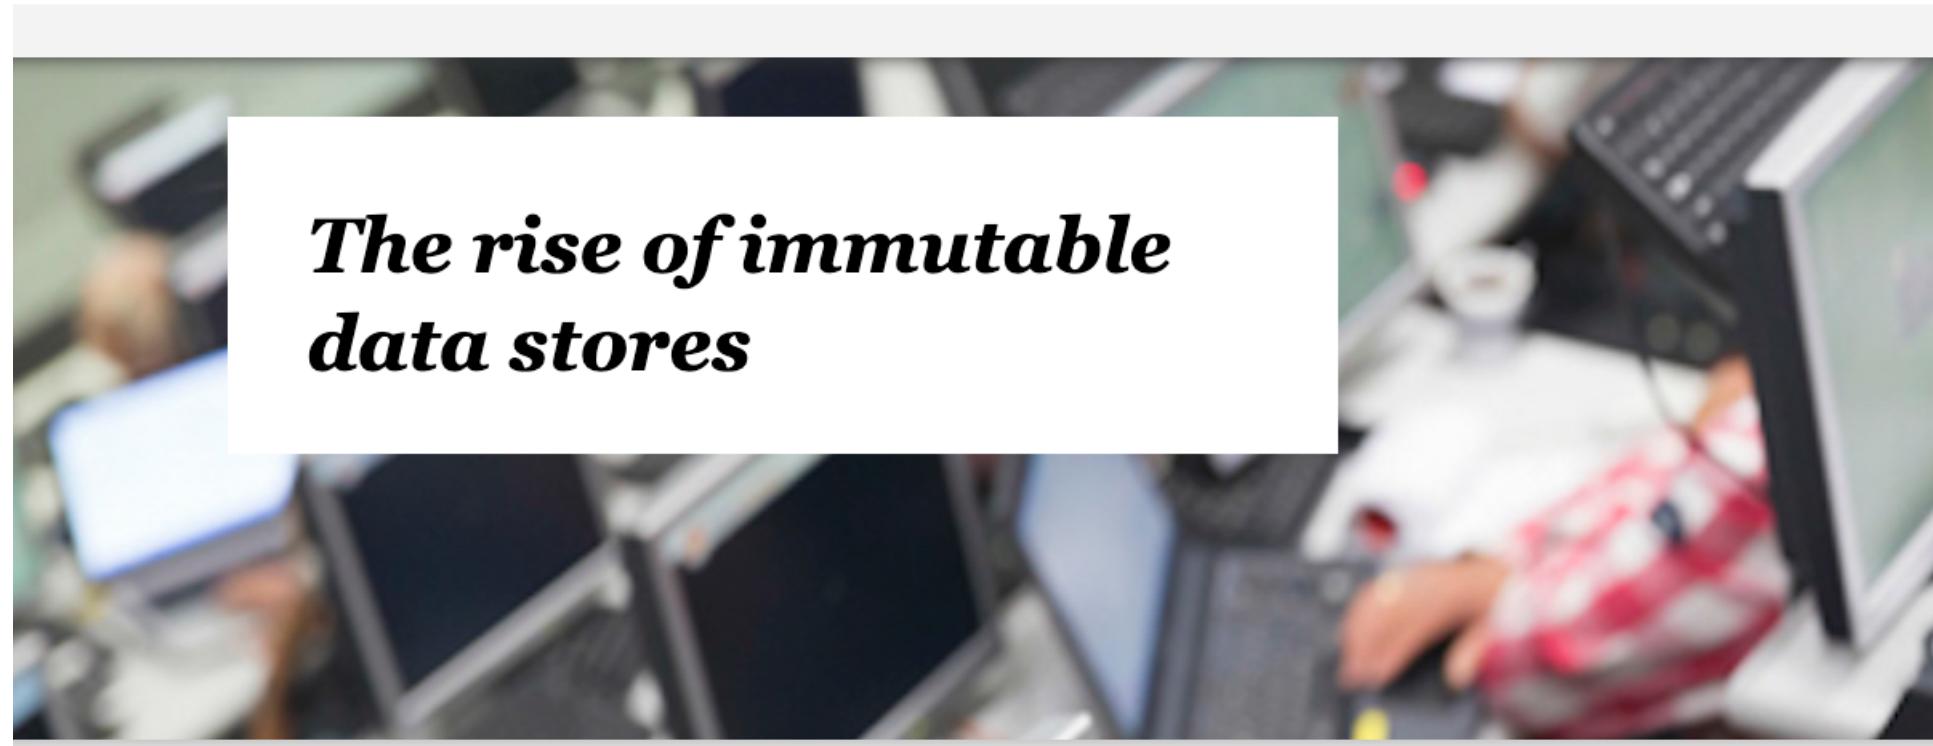
- » Haskell, Clojure, Erlang...
- » Immutable.js
- » CloudFormation, Terraform
- » Nix OS

**MY YOUTH  
DATA IS EPHERMAL**





# **ACCRETING FACTS OVER TIME**



## ***The rise of immutable data stores***

by [Alan Morrison](#)

September 8, 2015

---

Tags: [Data](#) [NoSQL Databases](#) [Unstructured Data](#)

### **Some innovators are abandoning long-held database principles. Why?**

The website for Room Key, a joint venture of six hotel chains to help travelers find and book lodging, collects data from as many as 17 million pages per month, records an average of 1.5 million events every 24 hours, and handles peak loads of 30 events per second. To process that onslaught of complex information, its database records each event without waiting for some other part of the system to do something first.

# REASONS WHY YOU SHOULD CARE

1. Immutable Database
2. Query Like a Ninja
3. Unbundle Your Database



# CRITIQUE OF SQL

- » lack of proper orthogonality
- » lack of compactness
- » lack of consistency
- » poor system cohesion

MAY 09, 2019

# We Can Do Better Than SQL

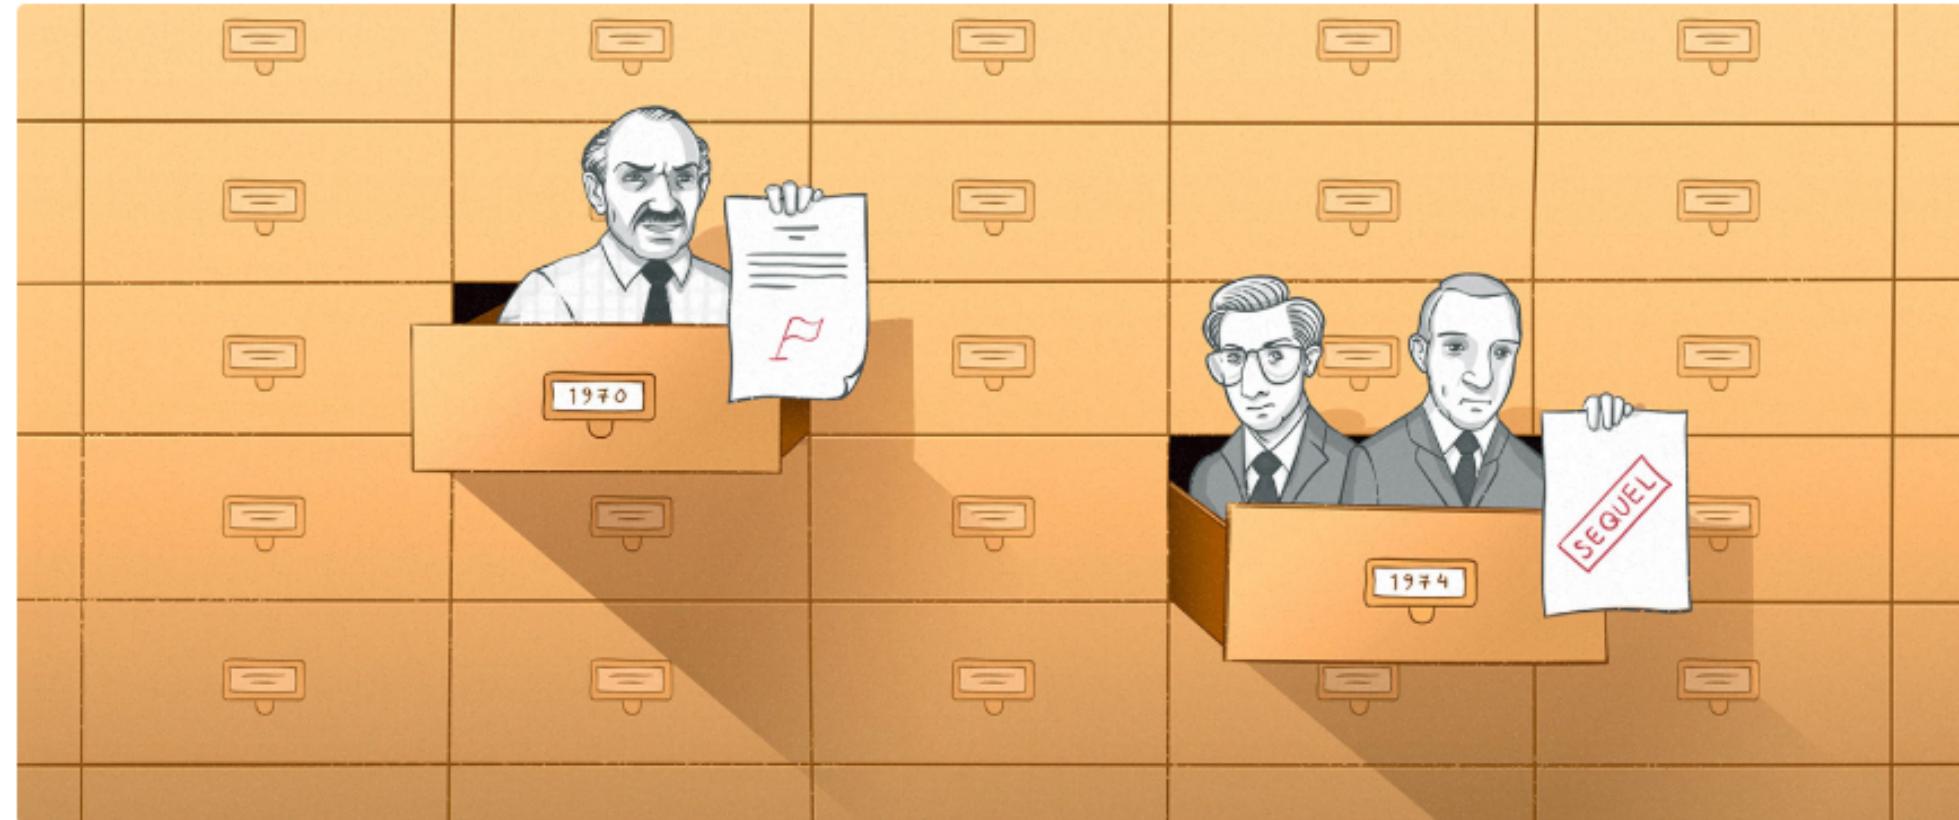
Elvis Pranskevichus

 [all blog posts](#)[A Brief History of the Relational Model and SQL](#)[Critique of SQL](#)[Why Does All This Matter?](#)[EdgeQL: Query Language Evolution](#)[Final Words](#)

tweet



subscribe



The questions we often hear are "Why create a new query language?" and "What's wrong with SQL?". This post contains answers to both.

Before we begin, let's overview some of the history of how the relational model came to be, and how SQL was created.

# **DATA ALOG**

## **(BOSS SLIDE)**

Datalog is a declarative query language with roots in logic programming that combines facts and rules to achieve the same power as relational algebra recursion.

```
SELECT greatest_songs.name
FROM artists
INNER JOIN greatest_songs
ON artists.id = greatest_songs.artist_id
WHERE artists.name = "Luan Santana";
```



```
[ :find ?s
:where
[?s :greatest-song/artist ?p]
[?p :artist/name "Luan Santana"]]
```

**HTTP://WWW.LEARNDATALOGTODAY.ORG**

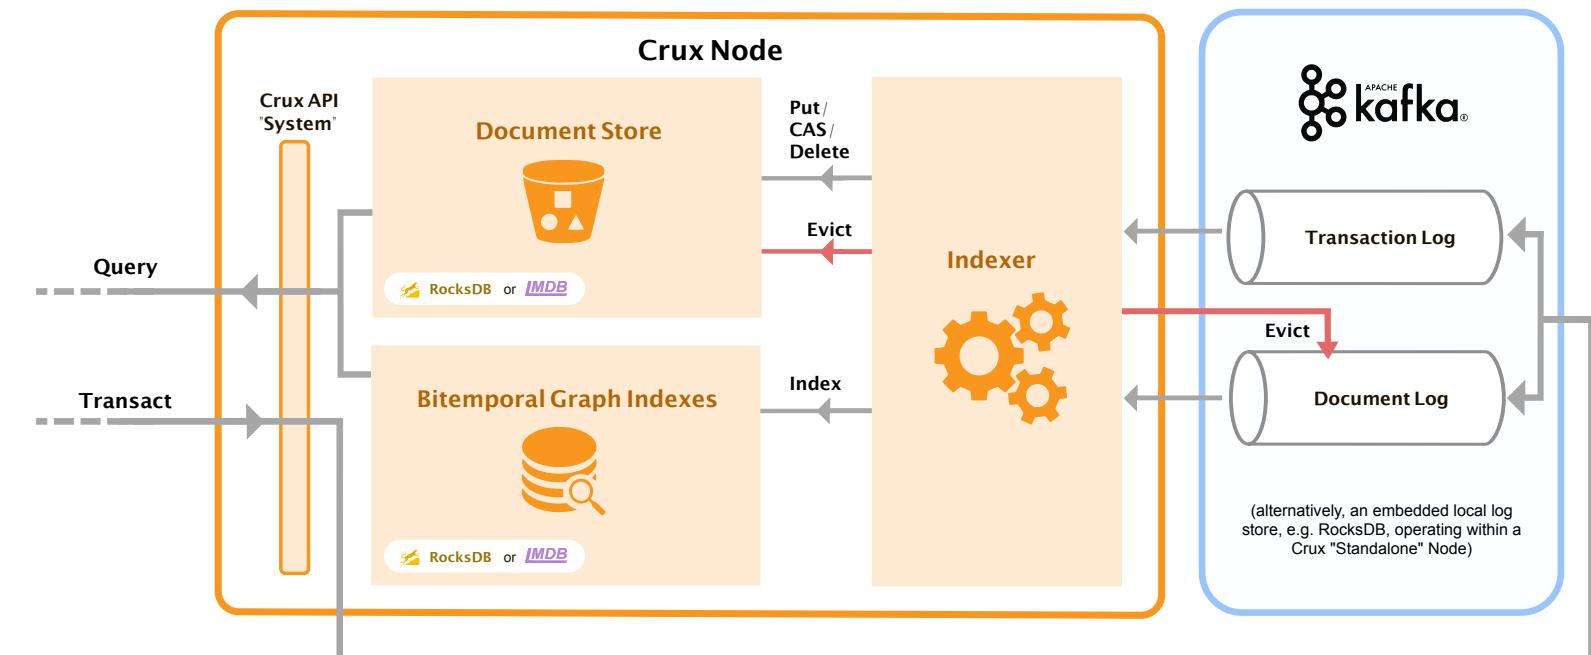
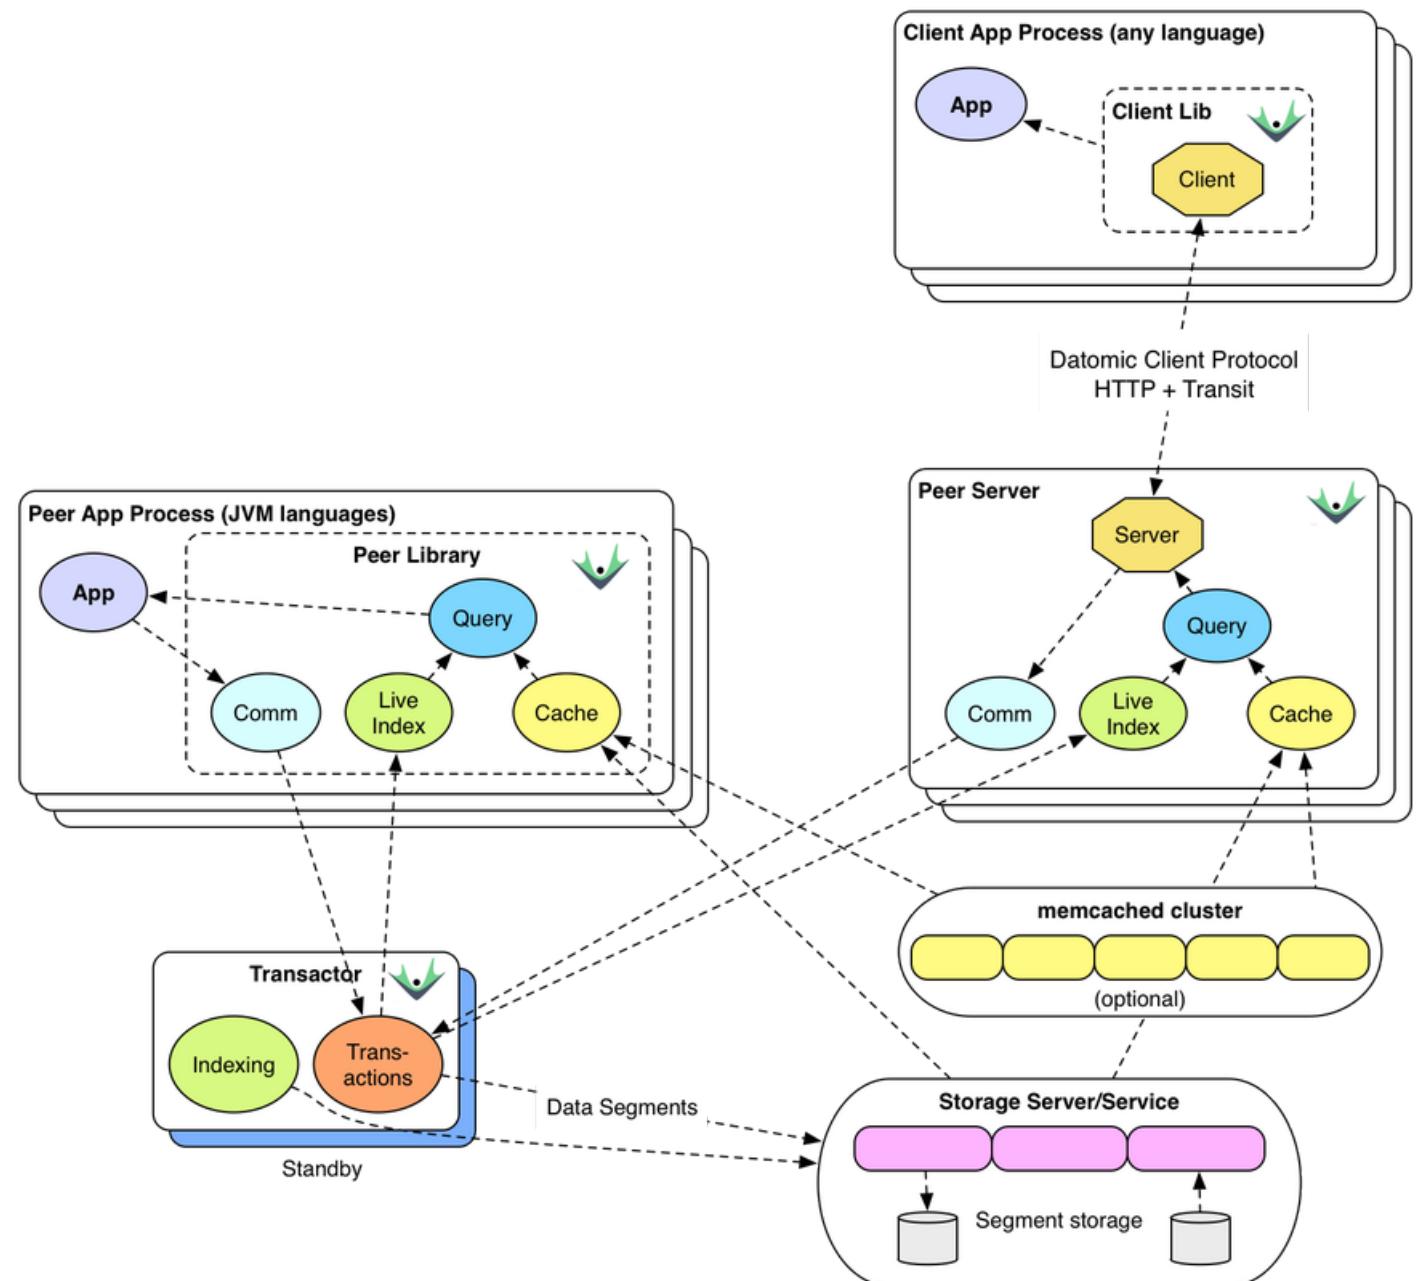
# REASONS WHY YOU SHOULD CARE

1. "Immutable" Database
2. Query Like a Zen Master
3. The Unbundled Database

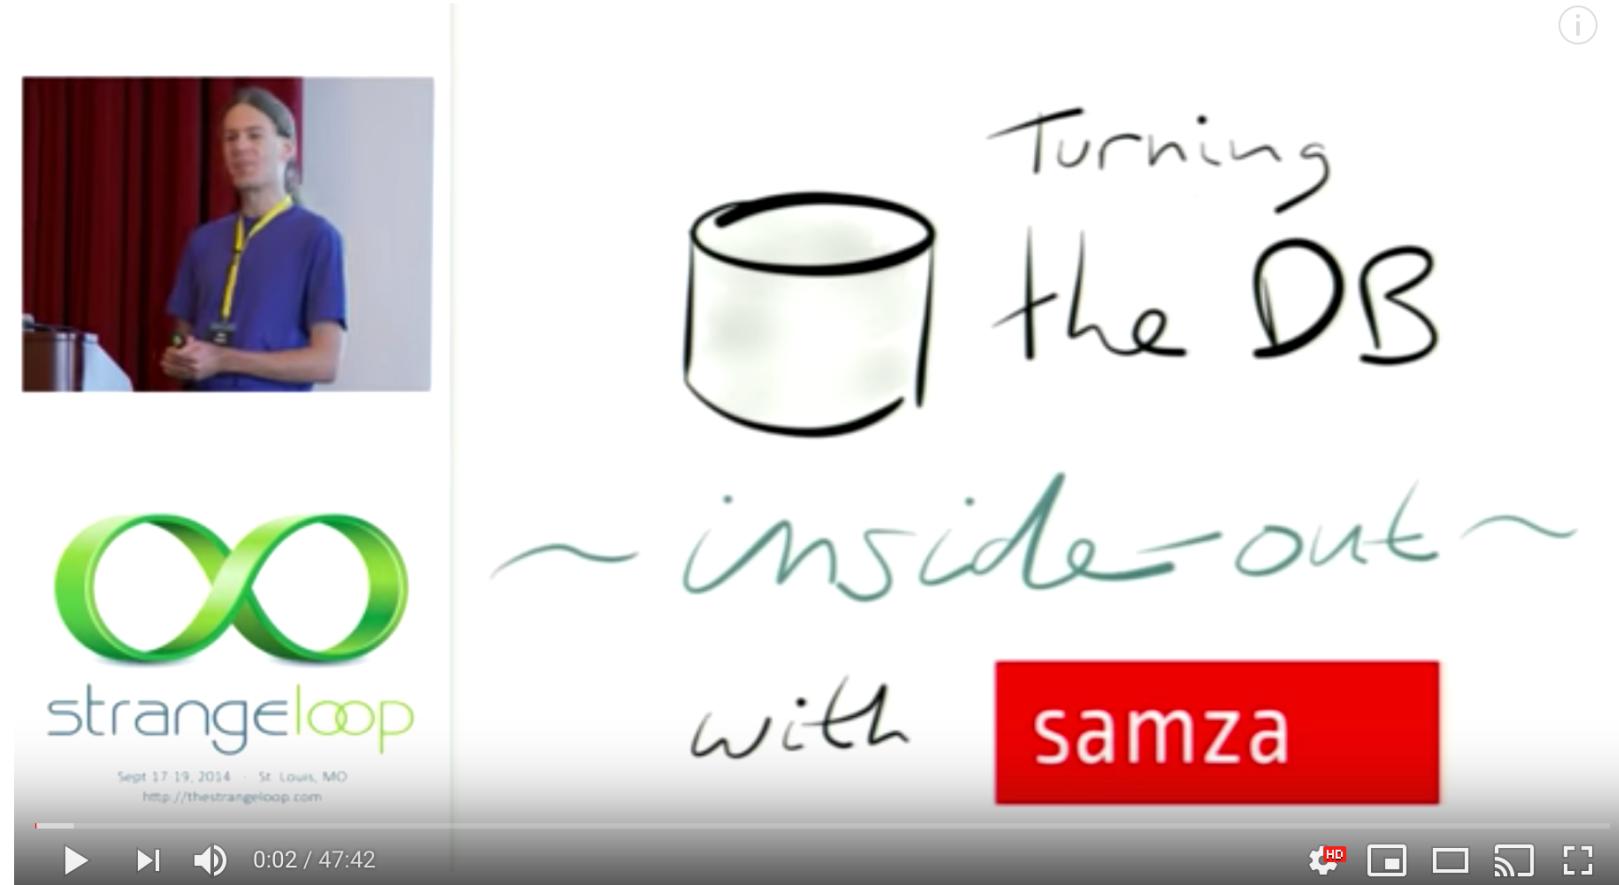
# THE UNBUNDLED DATABASE (BOSS SLIDE)

“What do we have to gain from turning the database inside out? Simpler code, better scalability, better robustness, lower latency, and more flexibility for doing interesting things with data.”

Martin Kleppmann



# "TURNING THE DATABASE INSIDE OUT"



<https://www.youtube.com/watch?v=fU9hR3kiOK0>

# **REASONS WHY YOU SHOULD CARE**

1. Immutable Database
2. Query Like a Ninja
3. Unbundle Your Database



C R U X



Datomic



# **THE 10 CRITICAL THINGS WHEN TALKING ABOUT DATOMIC AND CRUX**

# 1. LICENSING MODEL

# 1. LICENSING MODEL

- » Crux: Open Source (MIT)
- » Datomic: Proprietary Only

# **2. TEMPORAL MODEL**

## 2. TEMPORAL MODEL

- » Crux: Bitemporal
- » Datomic: Unitemporal

# DATOMIC'S TEMPORAL MODEL

- » "A datom is an immutable, point-in-time fact!"
- » Entity
- » Attribute
- » Value
- » Time
- » EAVT

<b>Entity</b>	<b>Attribute</b>	<b>Value</b>	<b>Time</b>
123	:person/likes	"pizza"	1978-07-08T20:19:03.176-00:00

<b>Entity</b>	<b>Attribute</b>	<b>Value</b>	<b>Time</b>
123	:person/likes	"pizza"	1978-07-08T20:19:03.176-00:00
123	:person/name	"John"	1978-07-08T20:19:03.176-00:00
123	:contact/phone	"+1 (305) 555-5524"	2009-03-01T14:24:32.074-00:00
123	:contact/phone	"+1 (646) 341-3367"	2017-10-25T08:22:06.165-00:00

# CRUX'S BITEMPORAL MODEL

- » Transaction Time
- » Valid Time









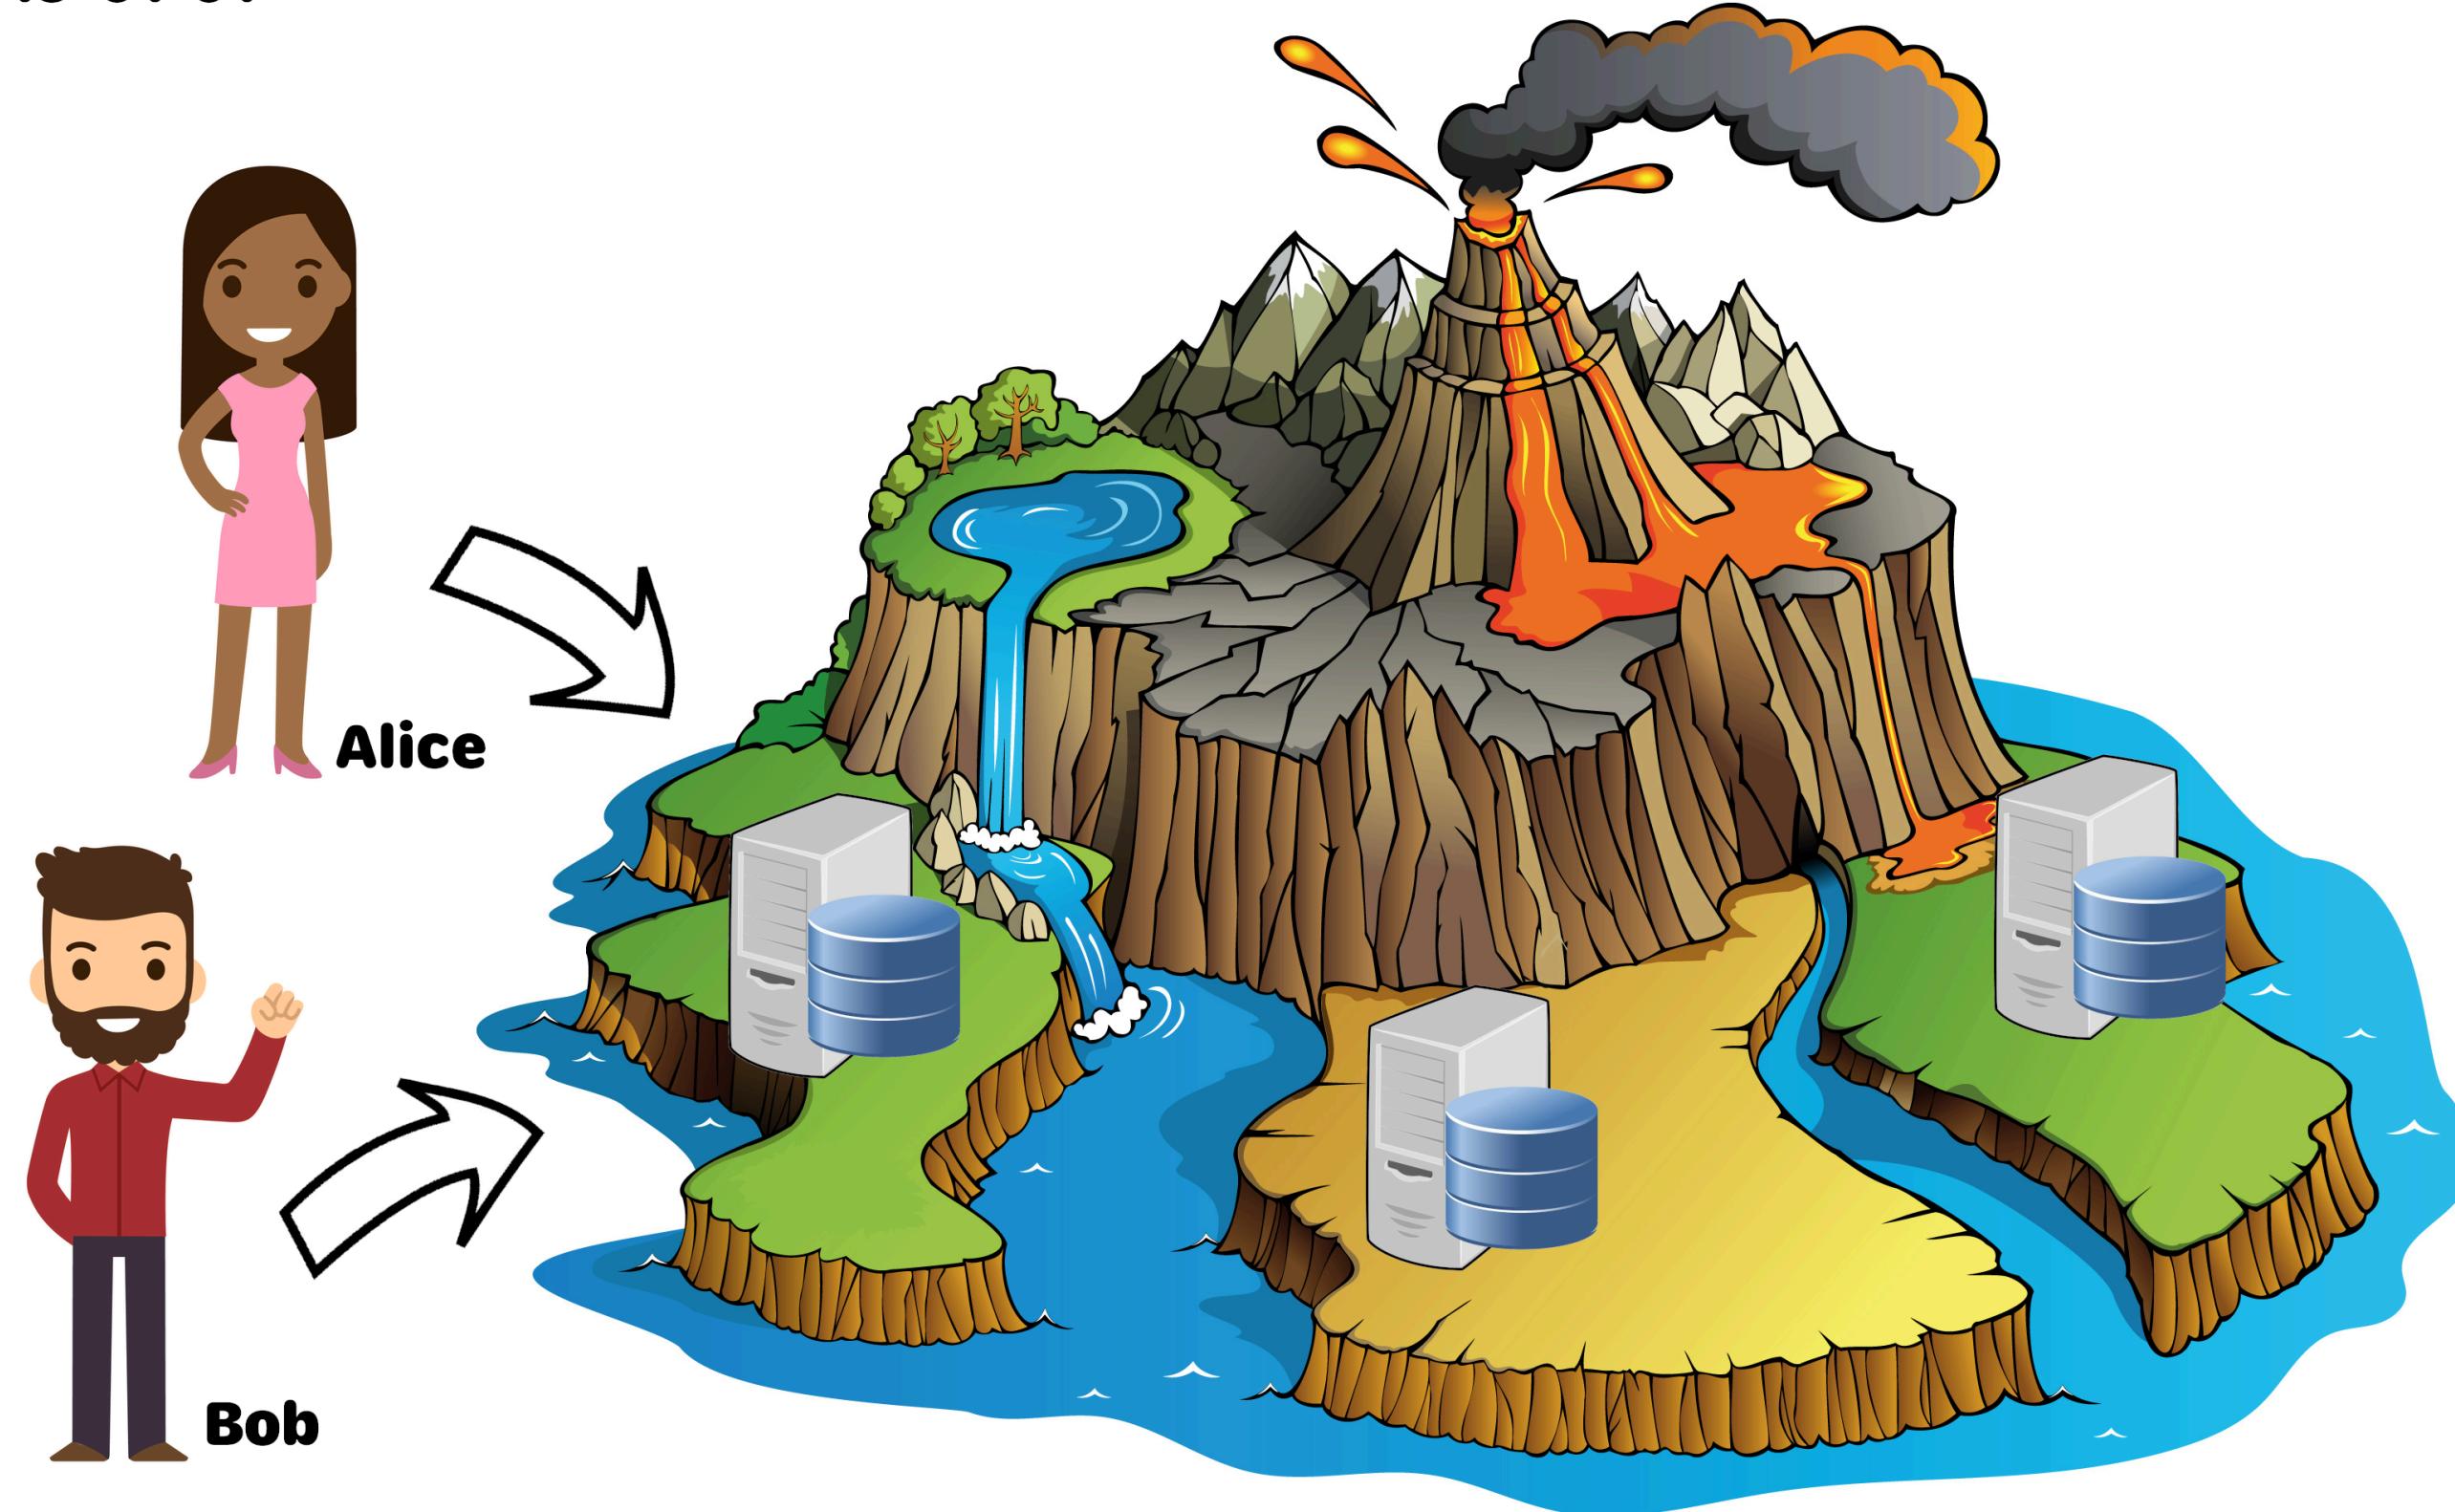




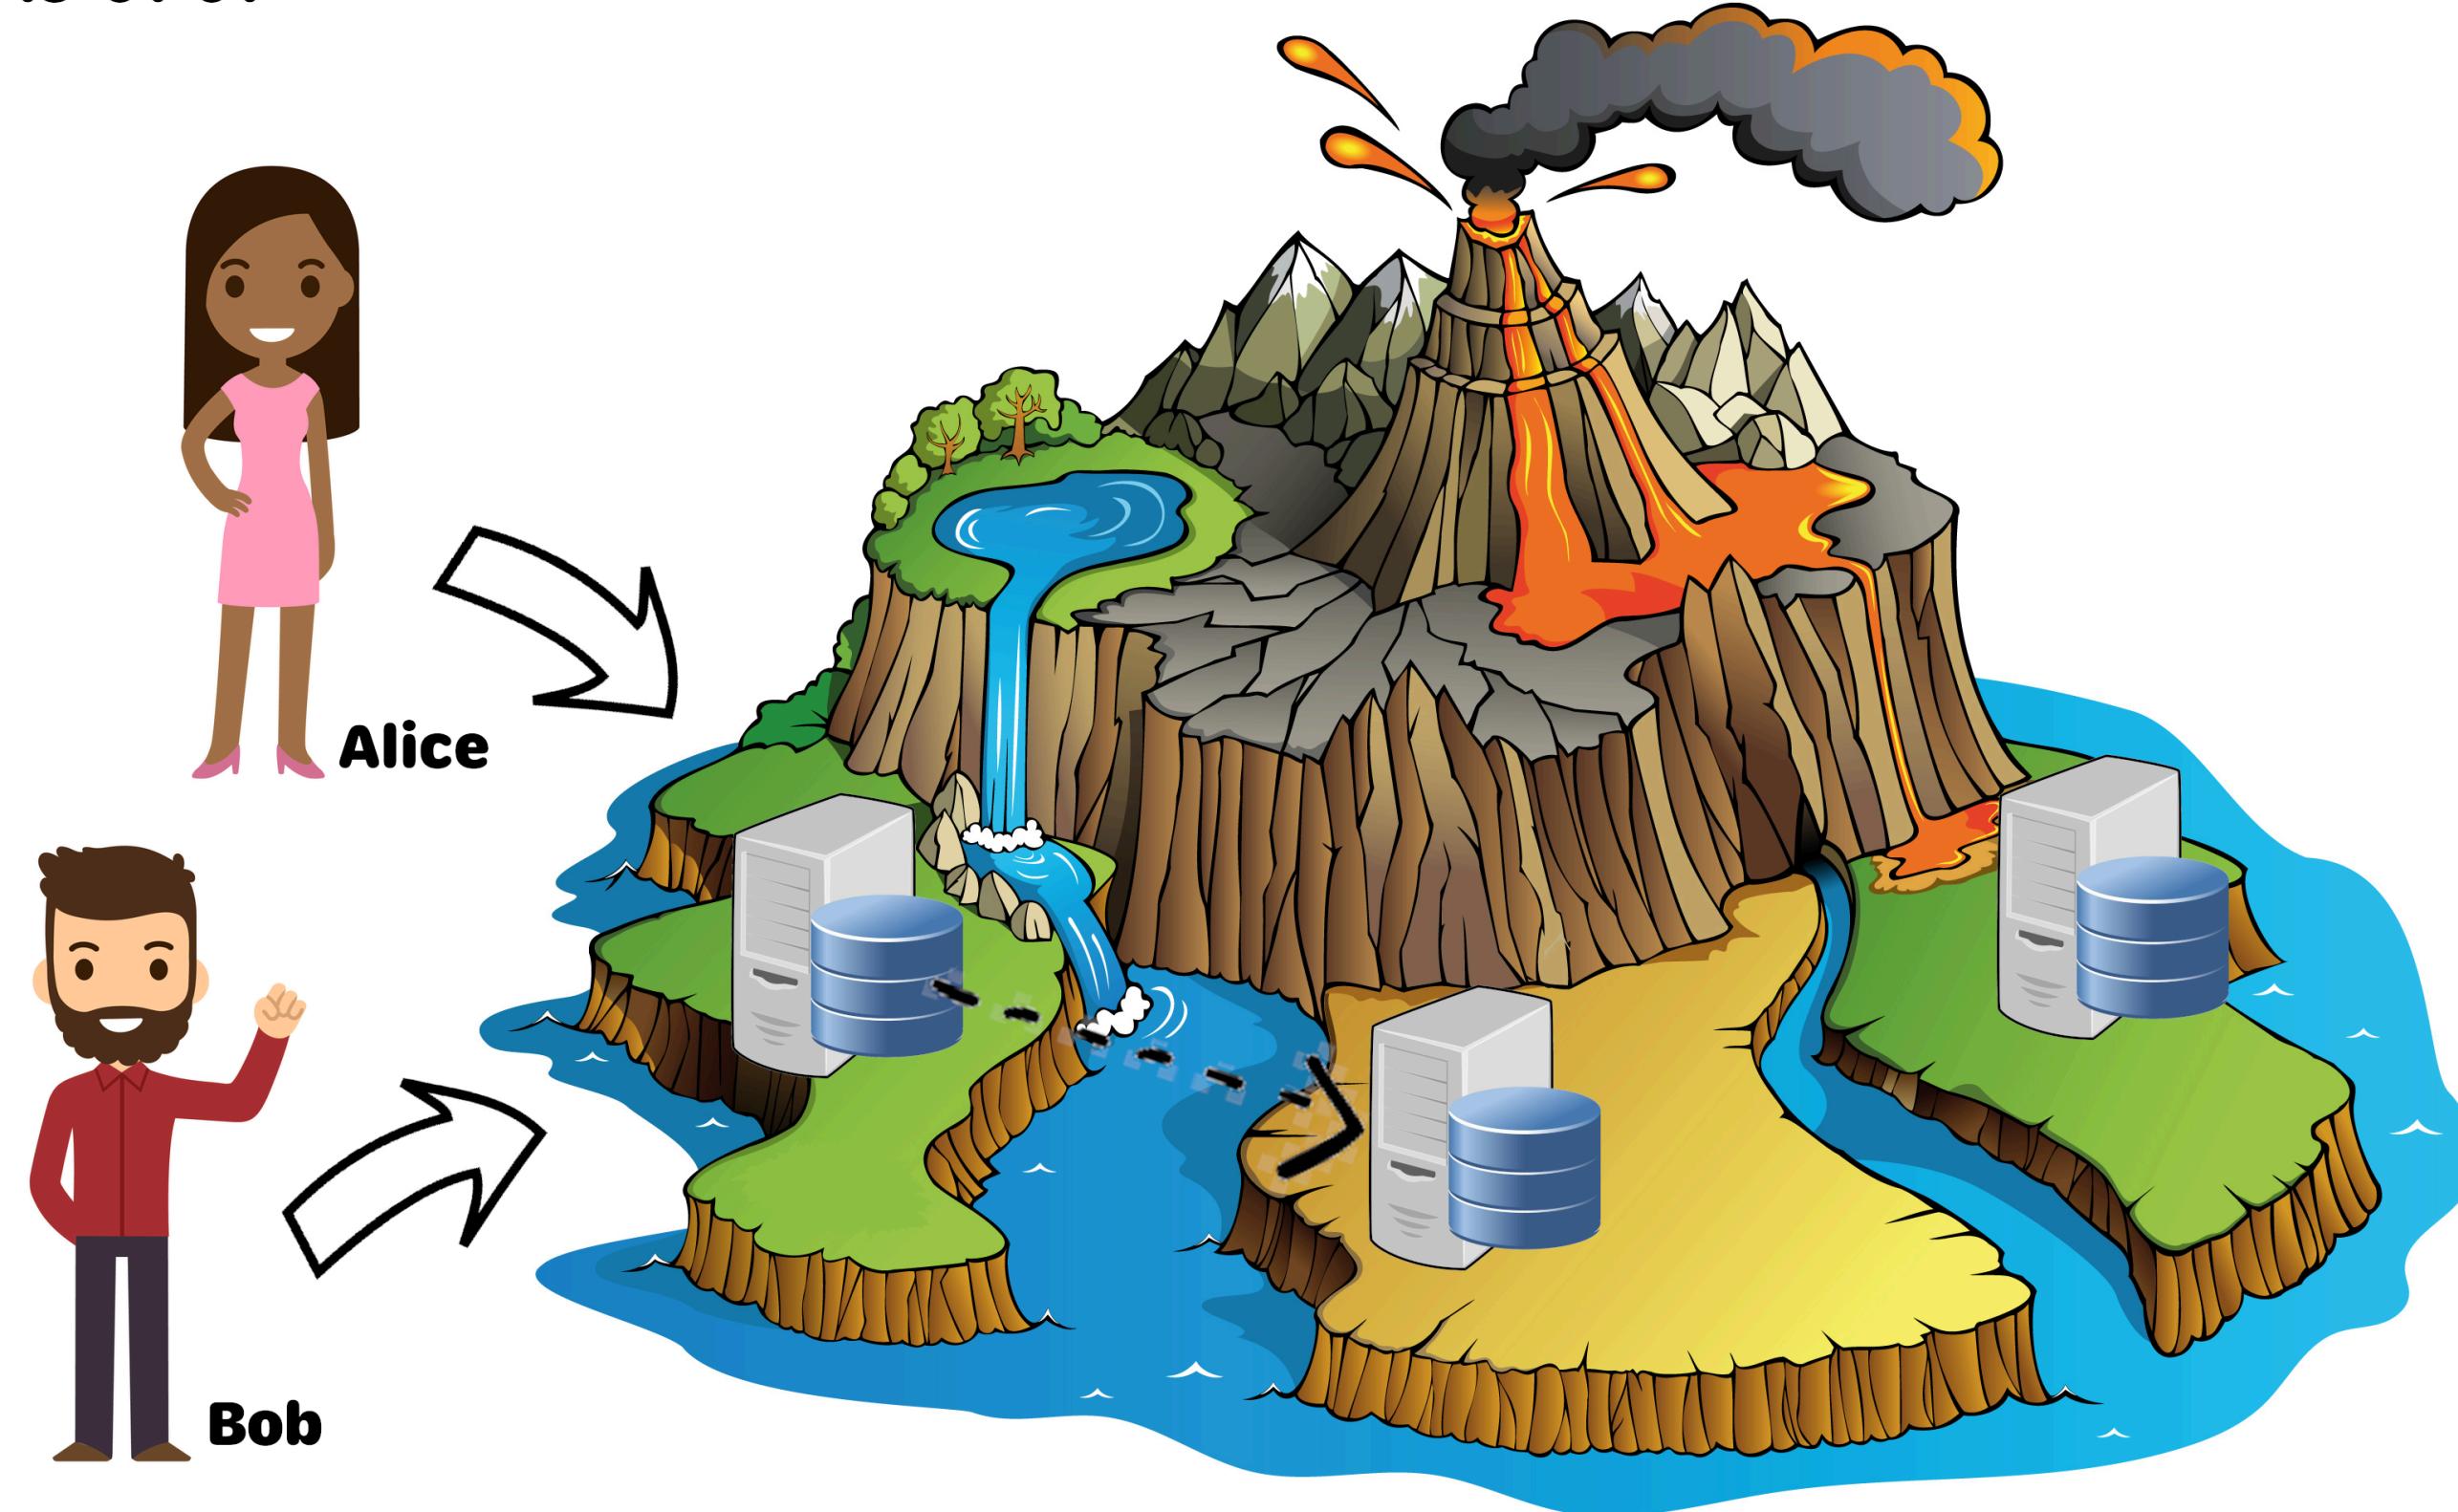
2019-01-01



2019-01-01



2019-01-01



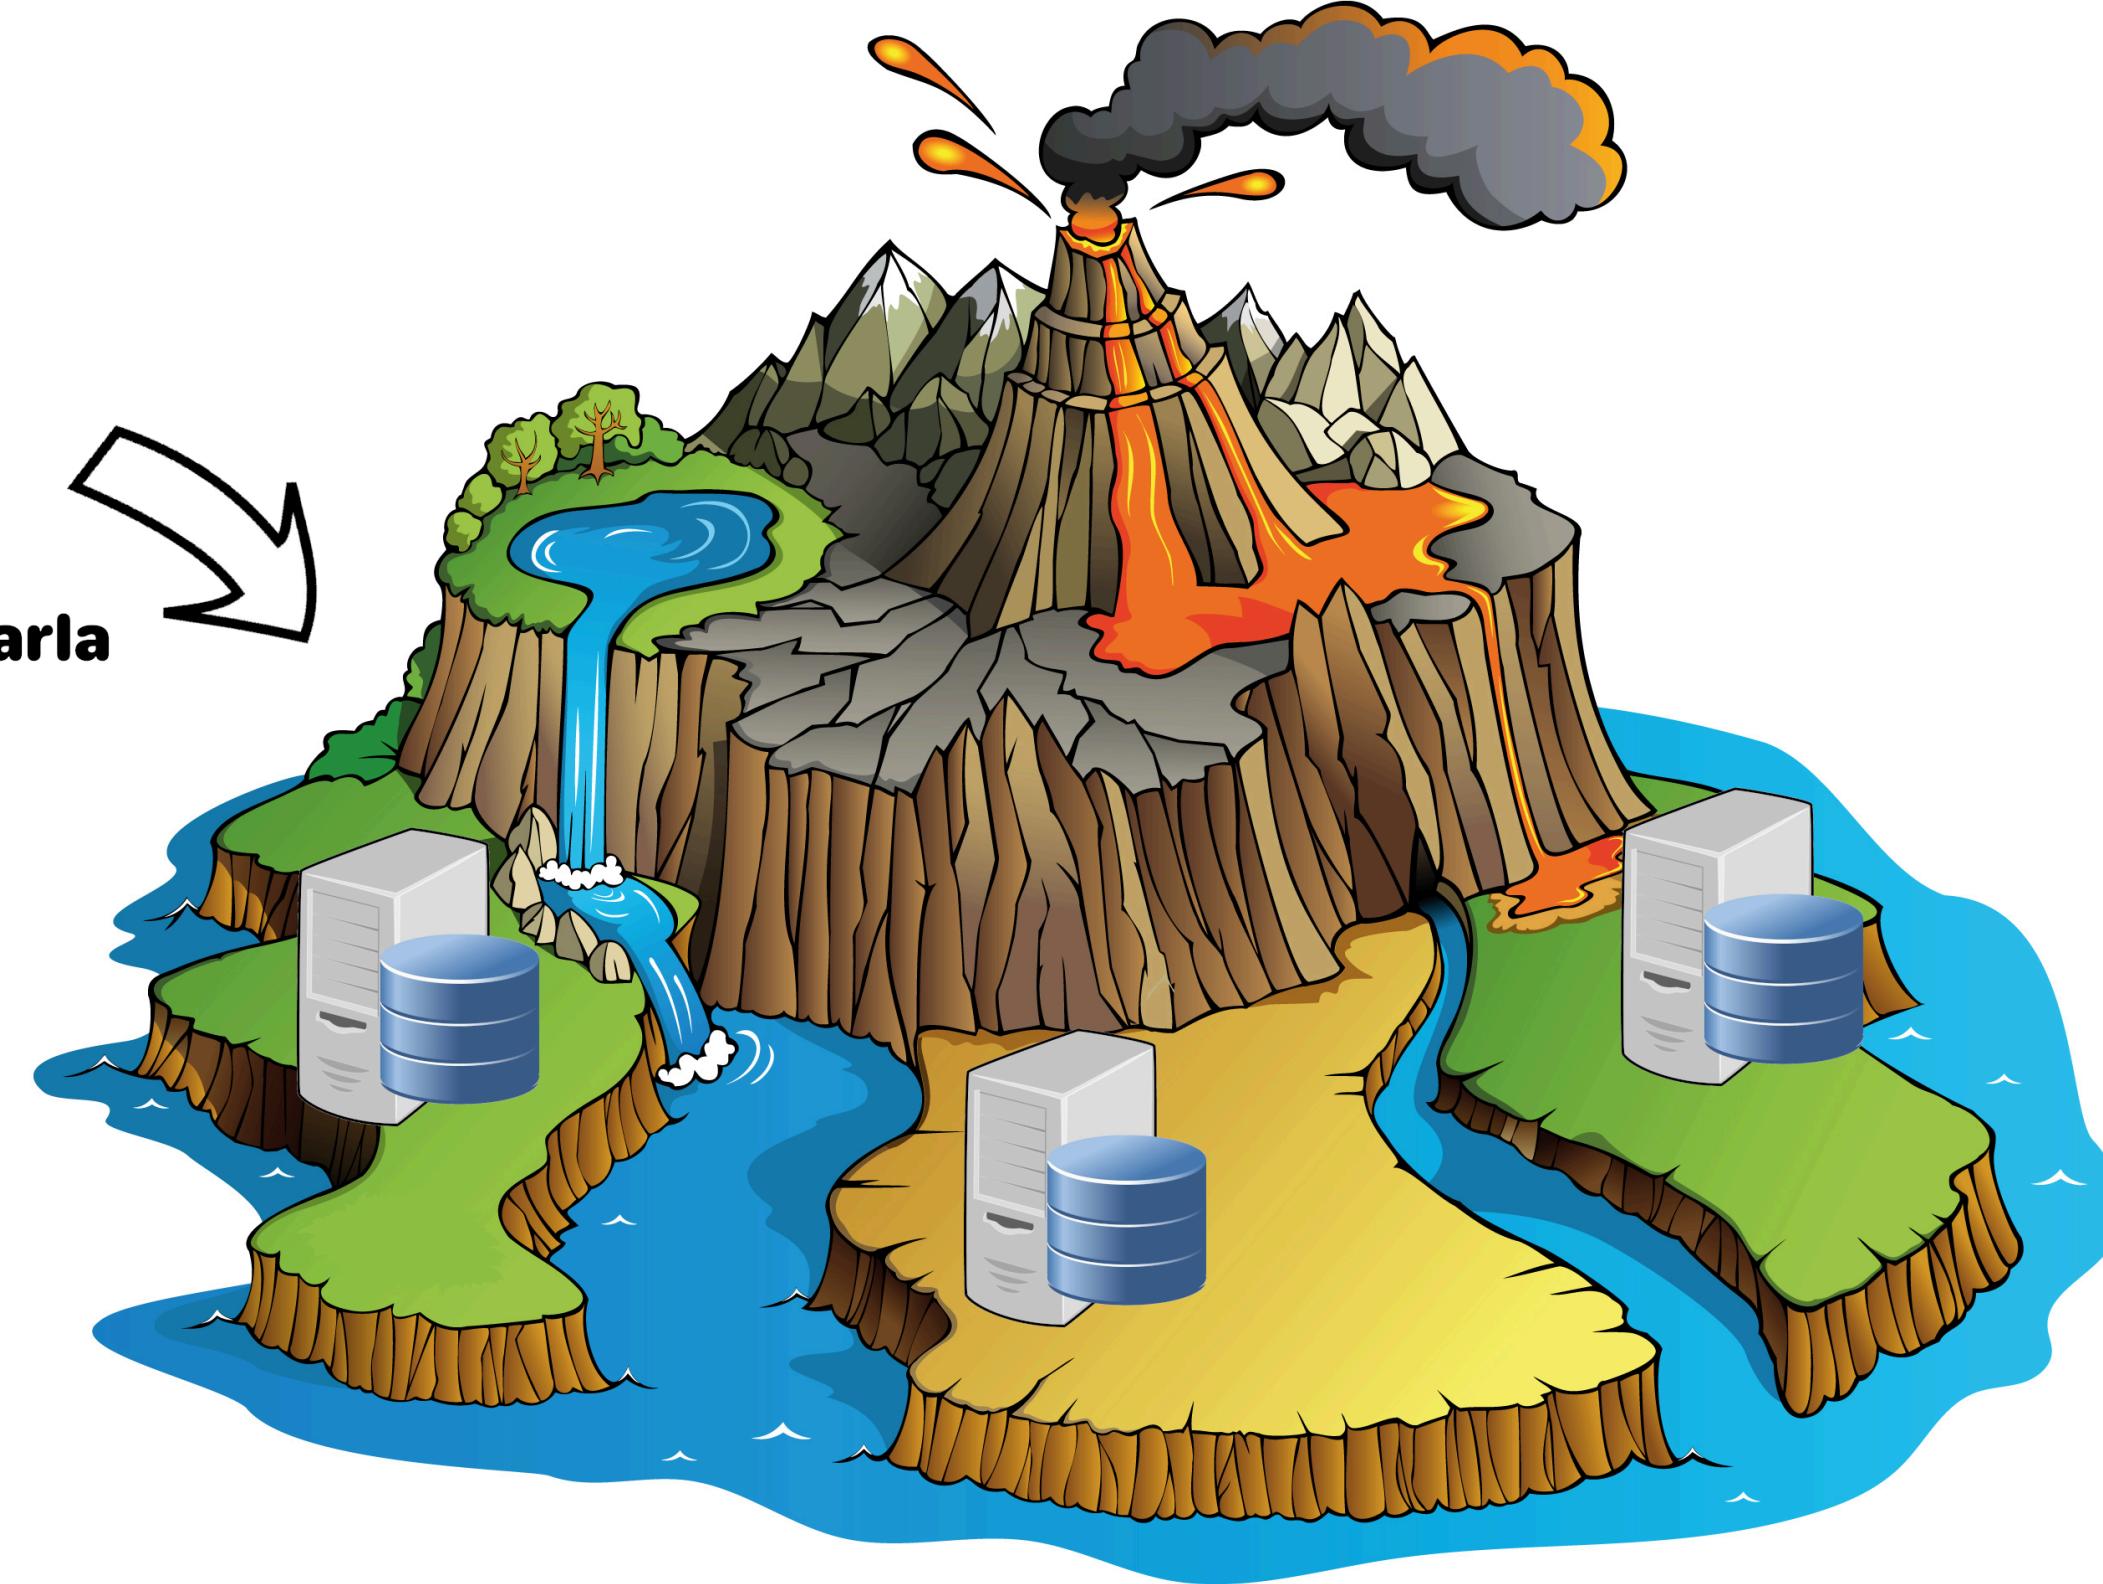
**2019-01-02**



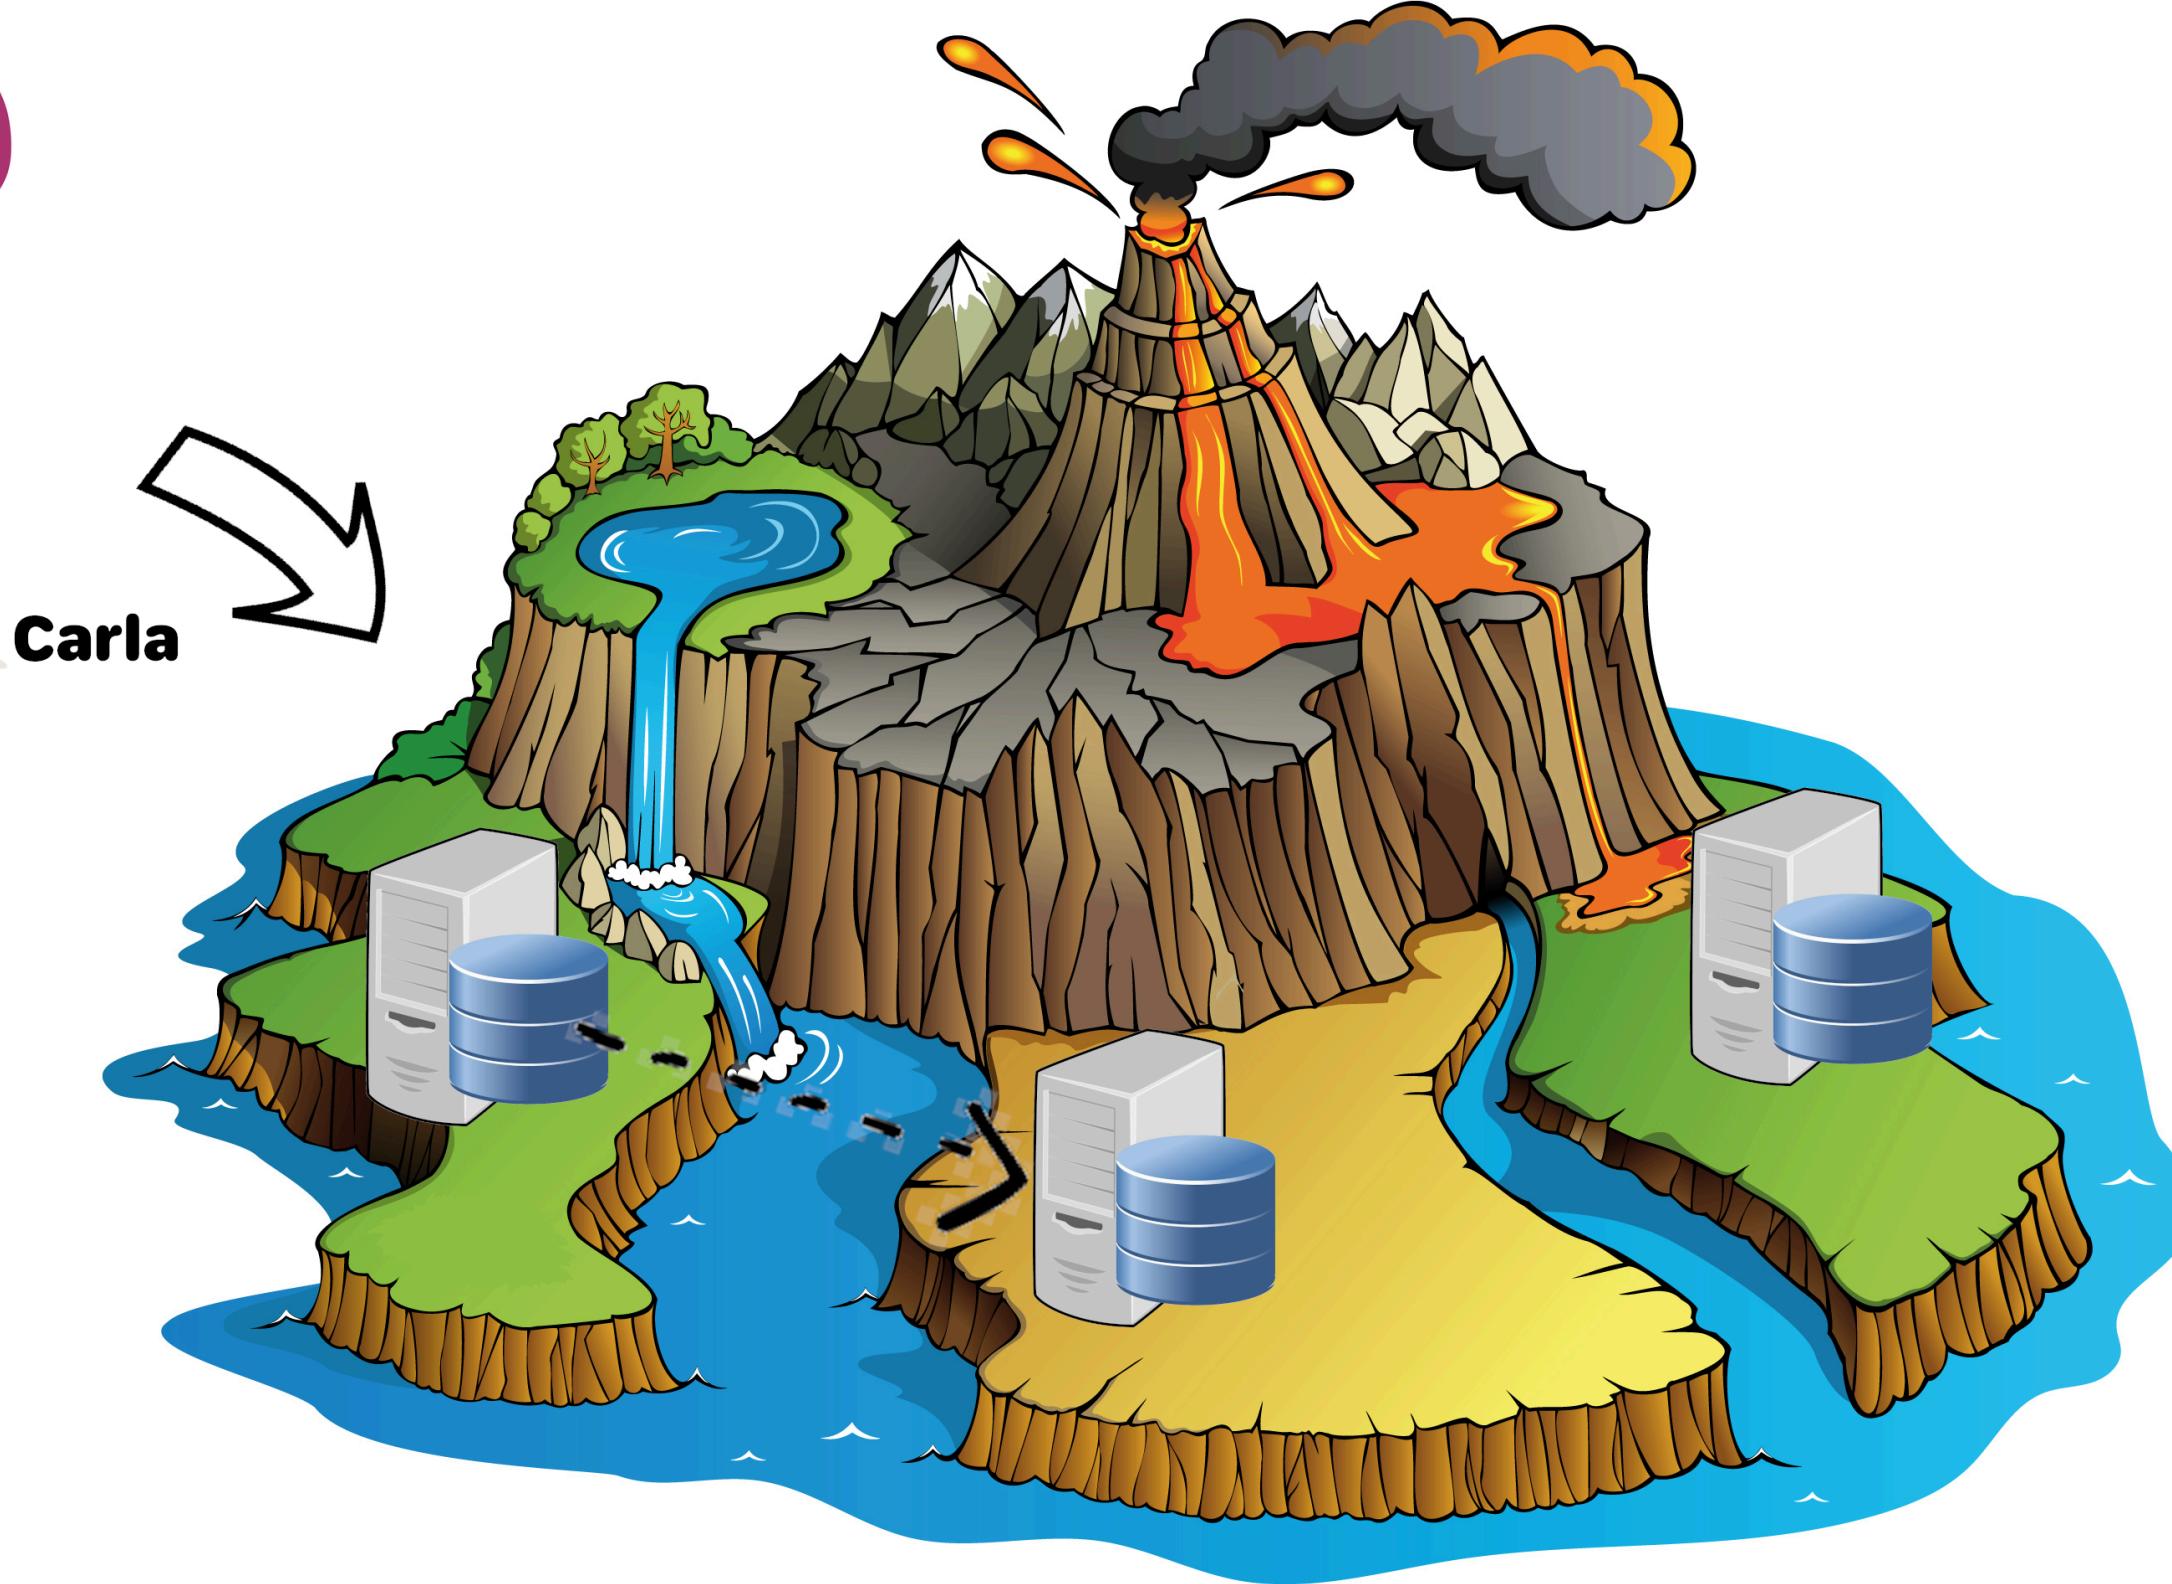
2019-01-02



Carla



2019-01-02



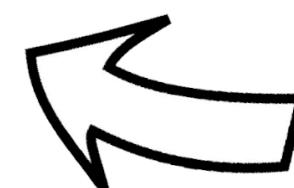
**2019-01-03**



2019-01-03



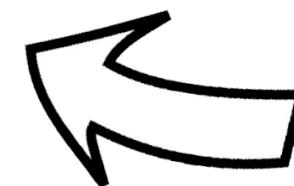
Carla



2019-01-03



Carla



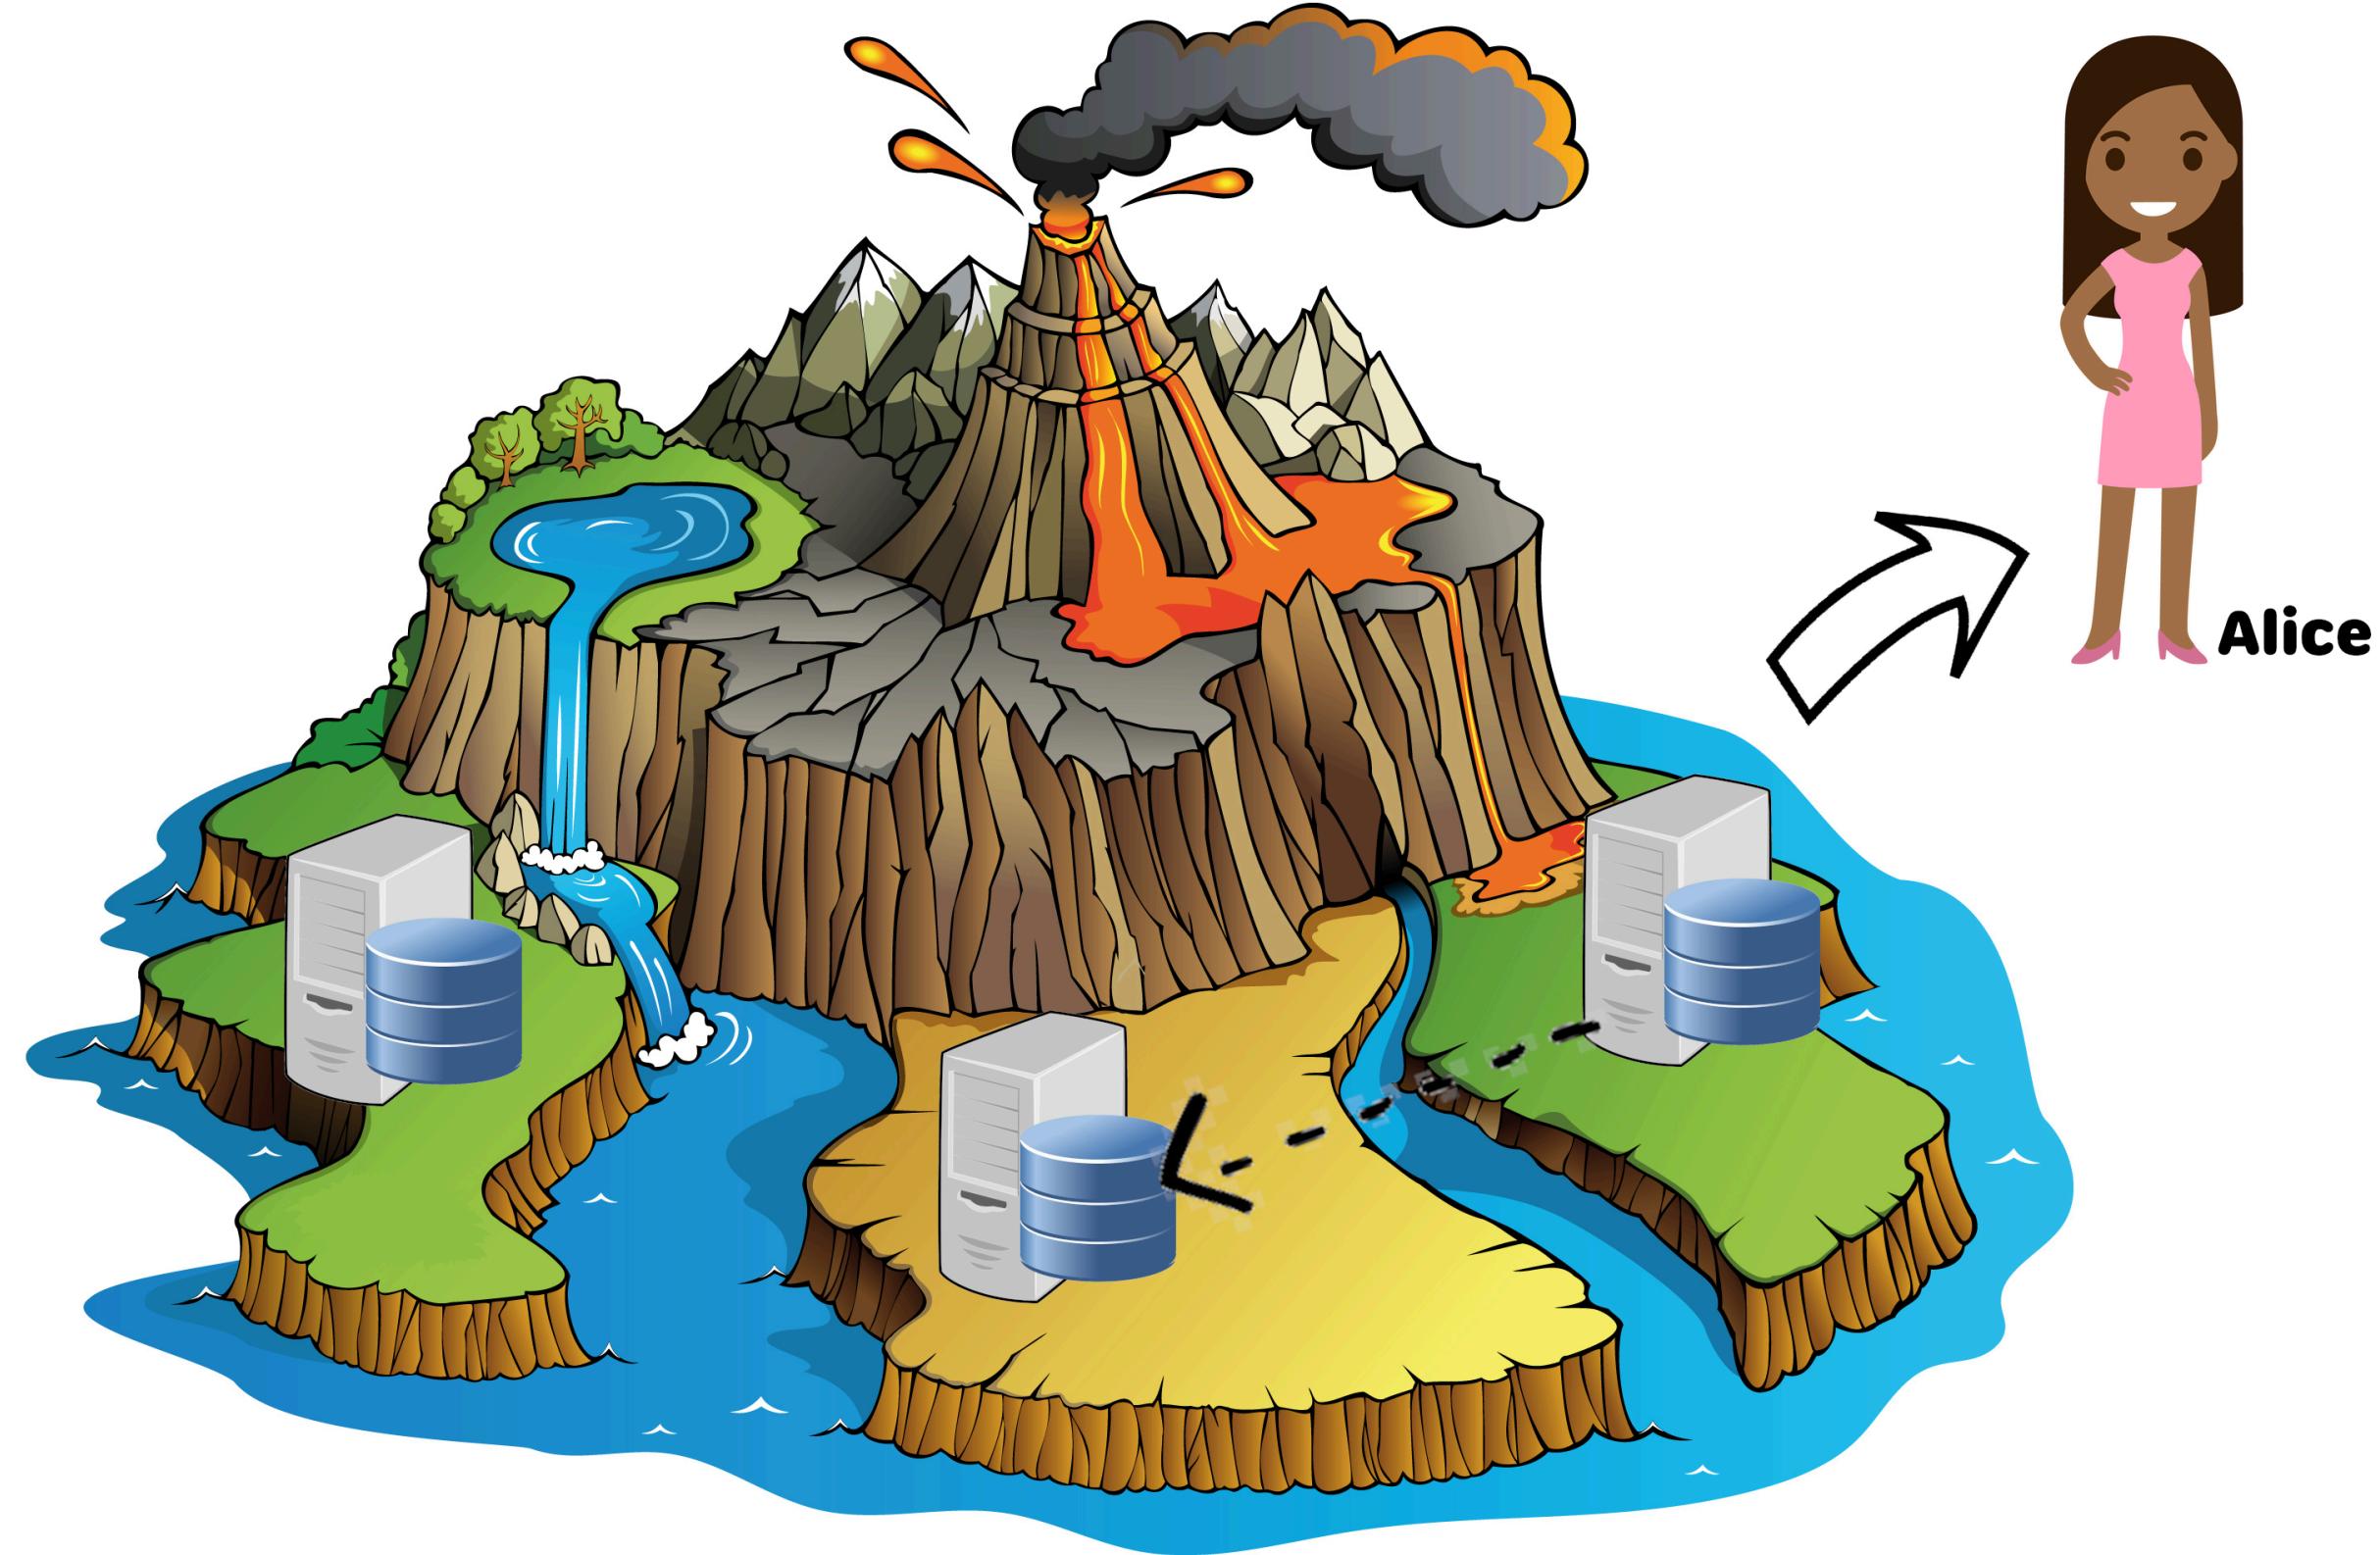
**2019-01-04**



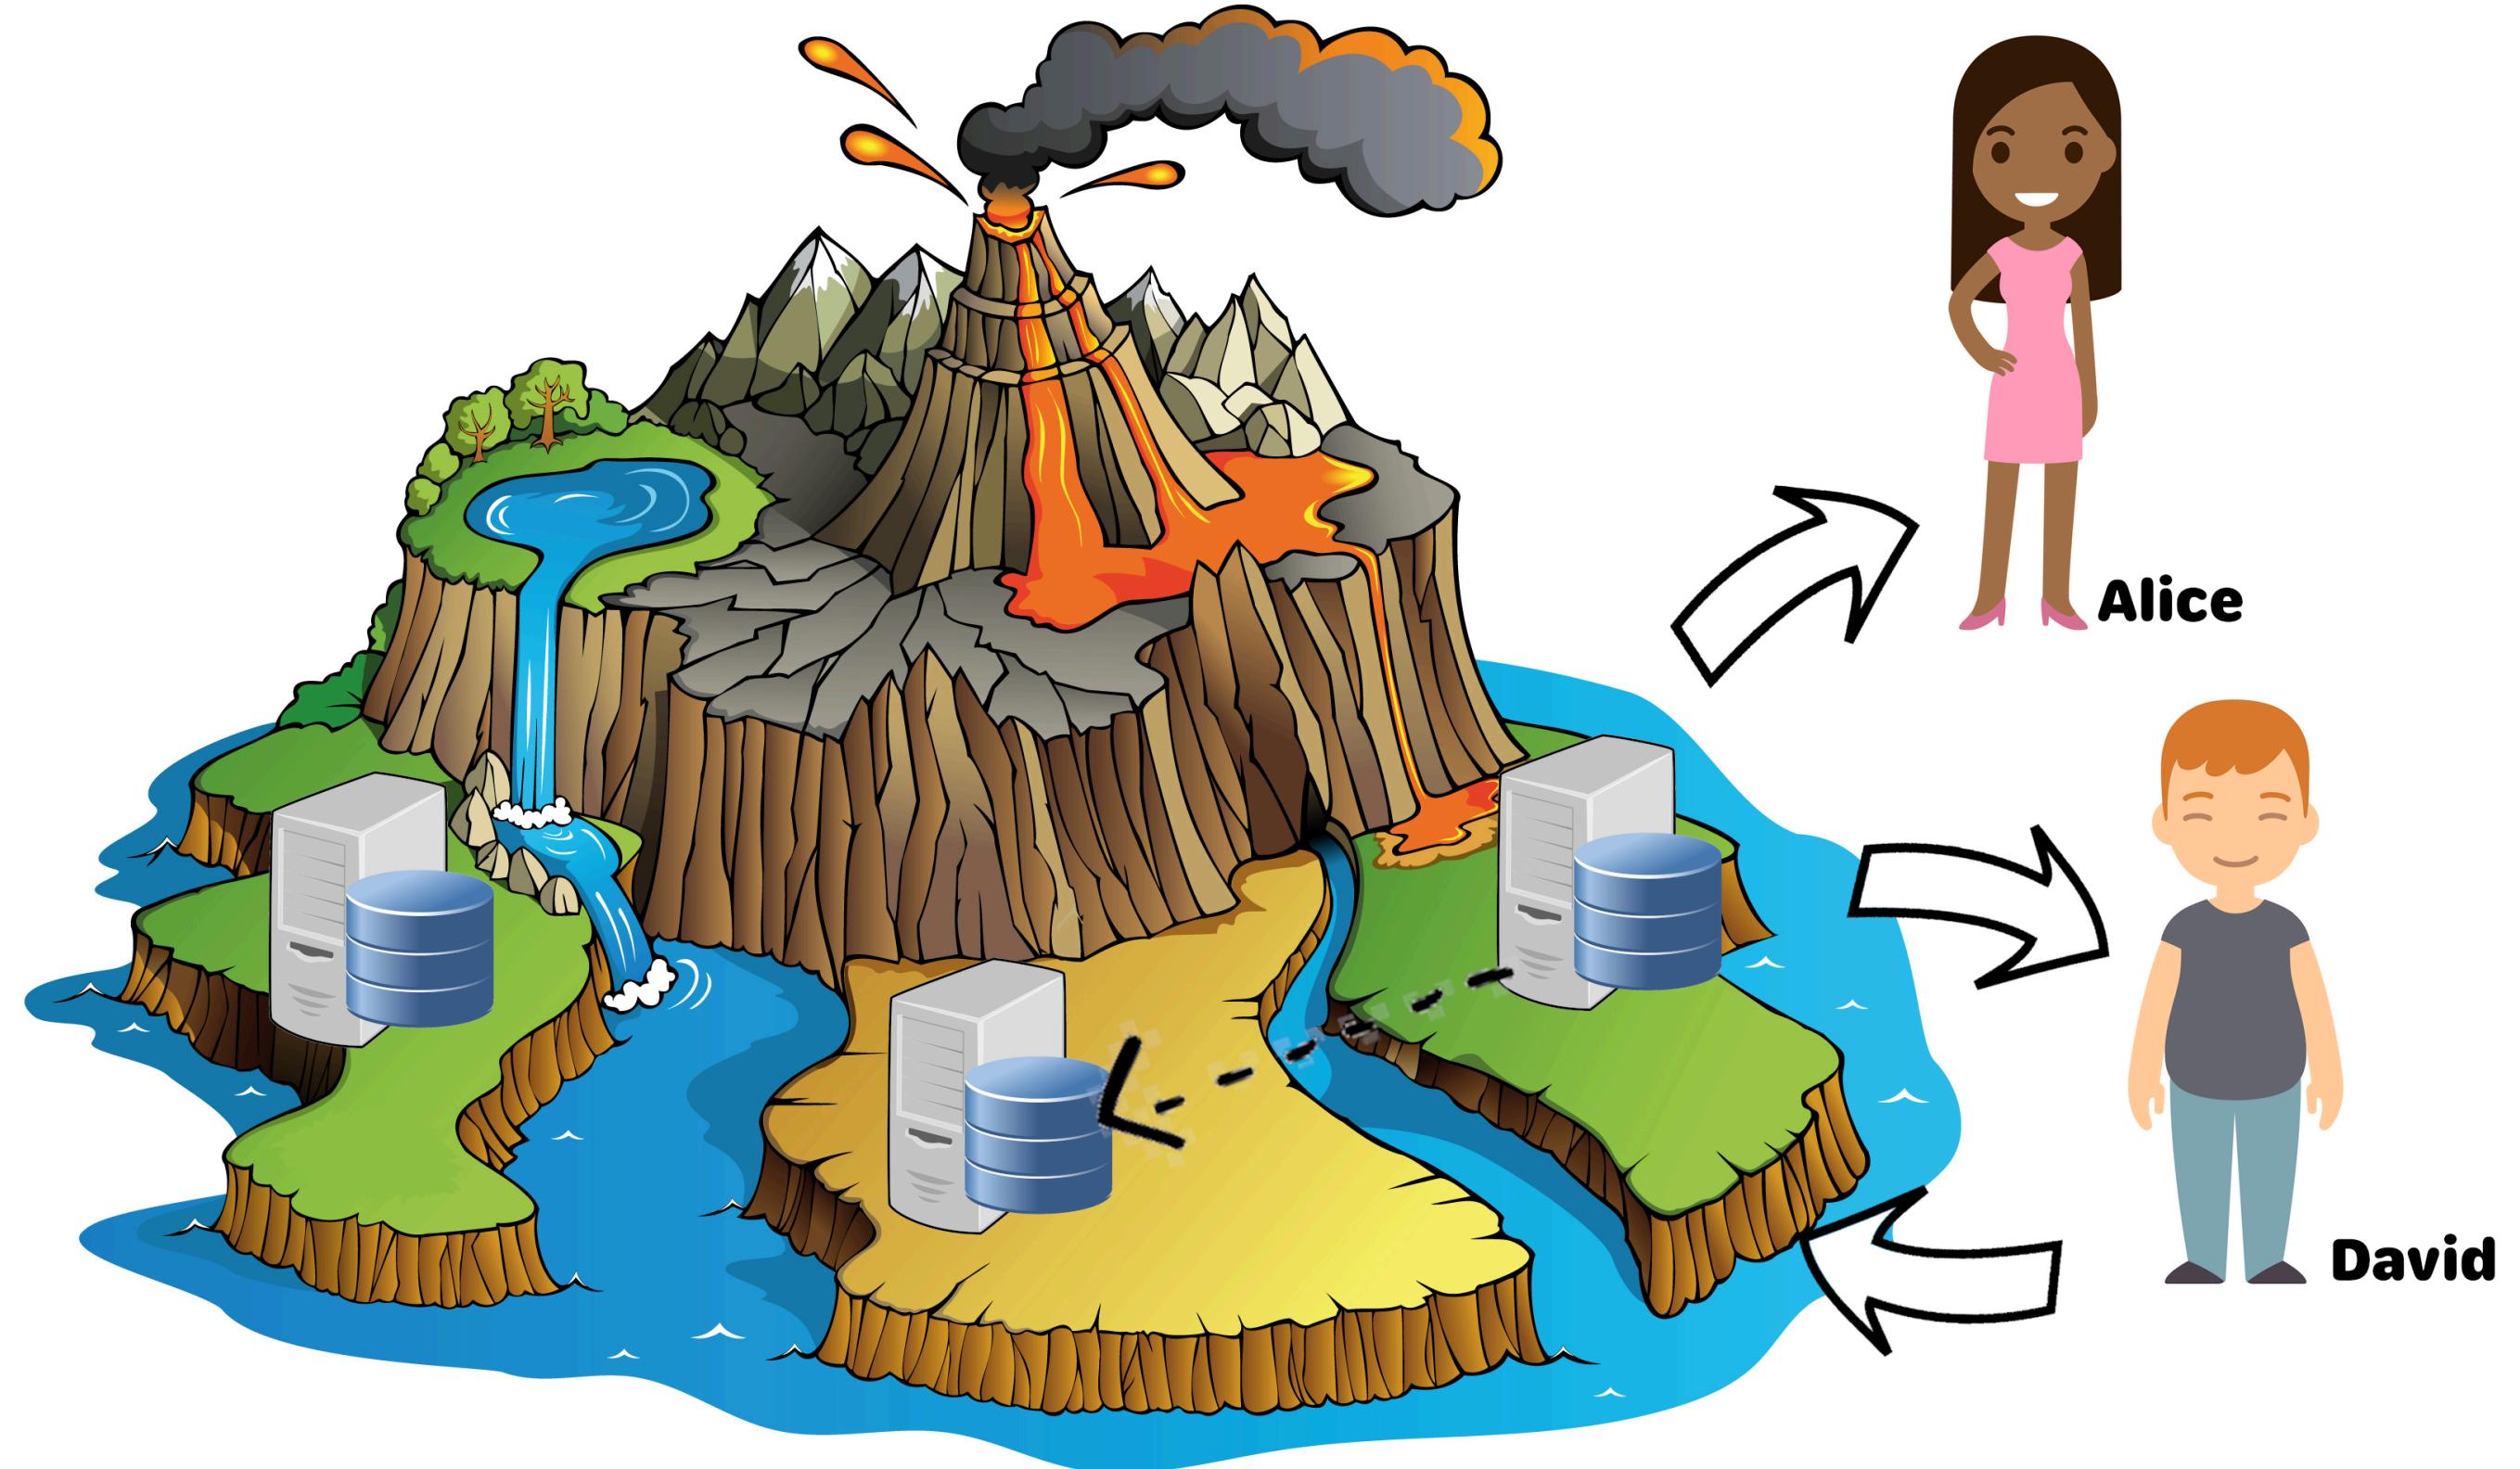
**2019-01-04**



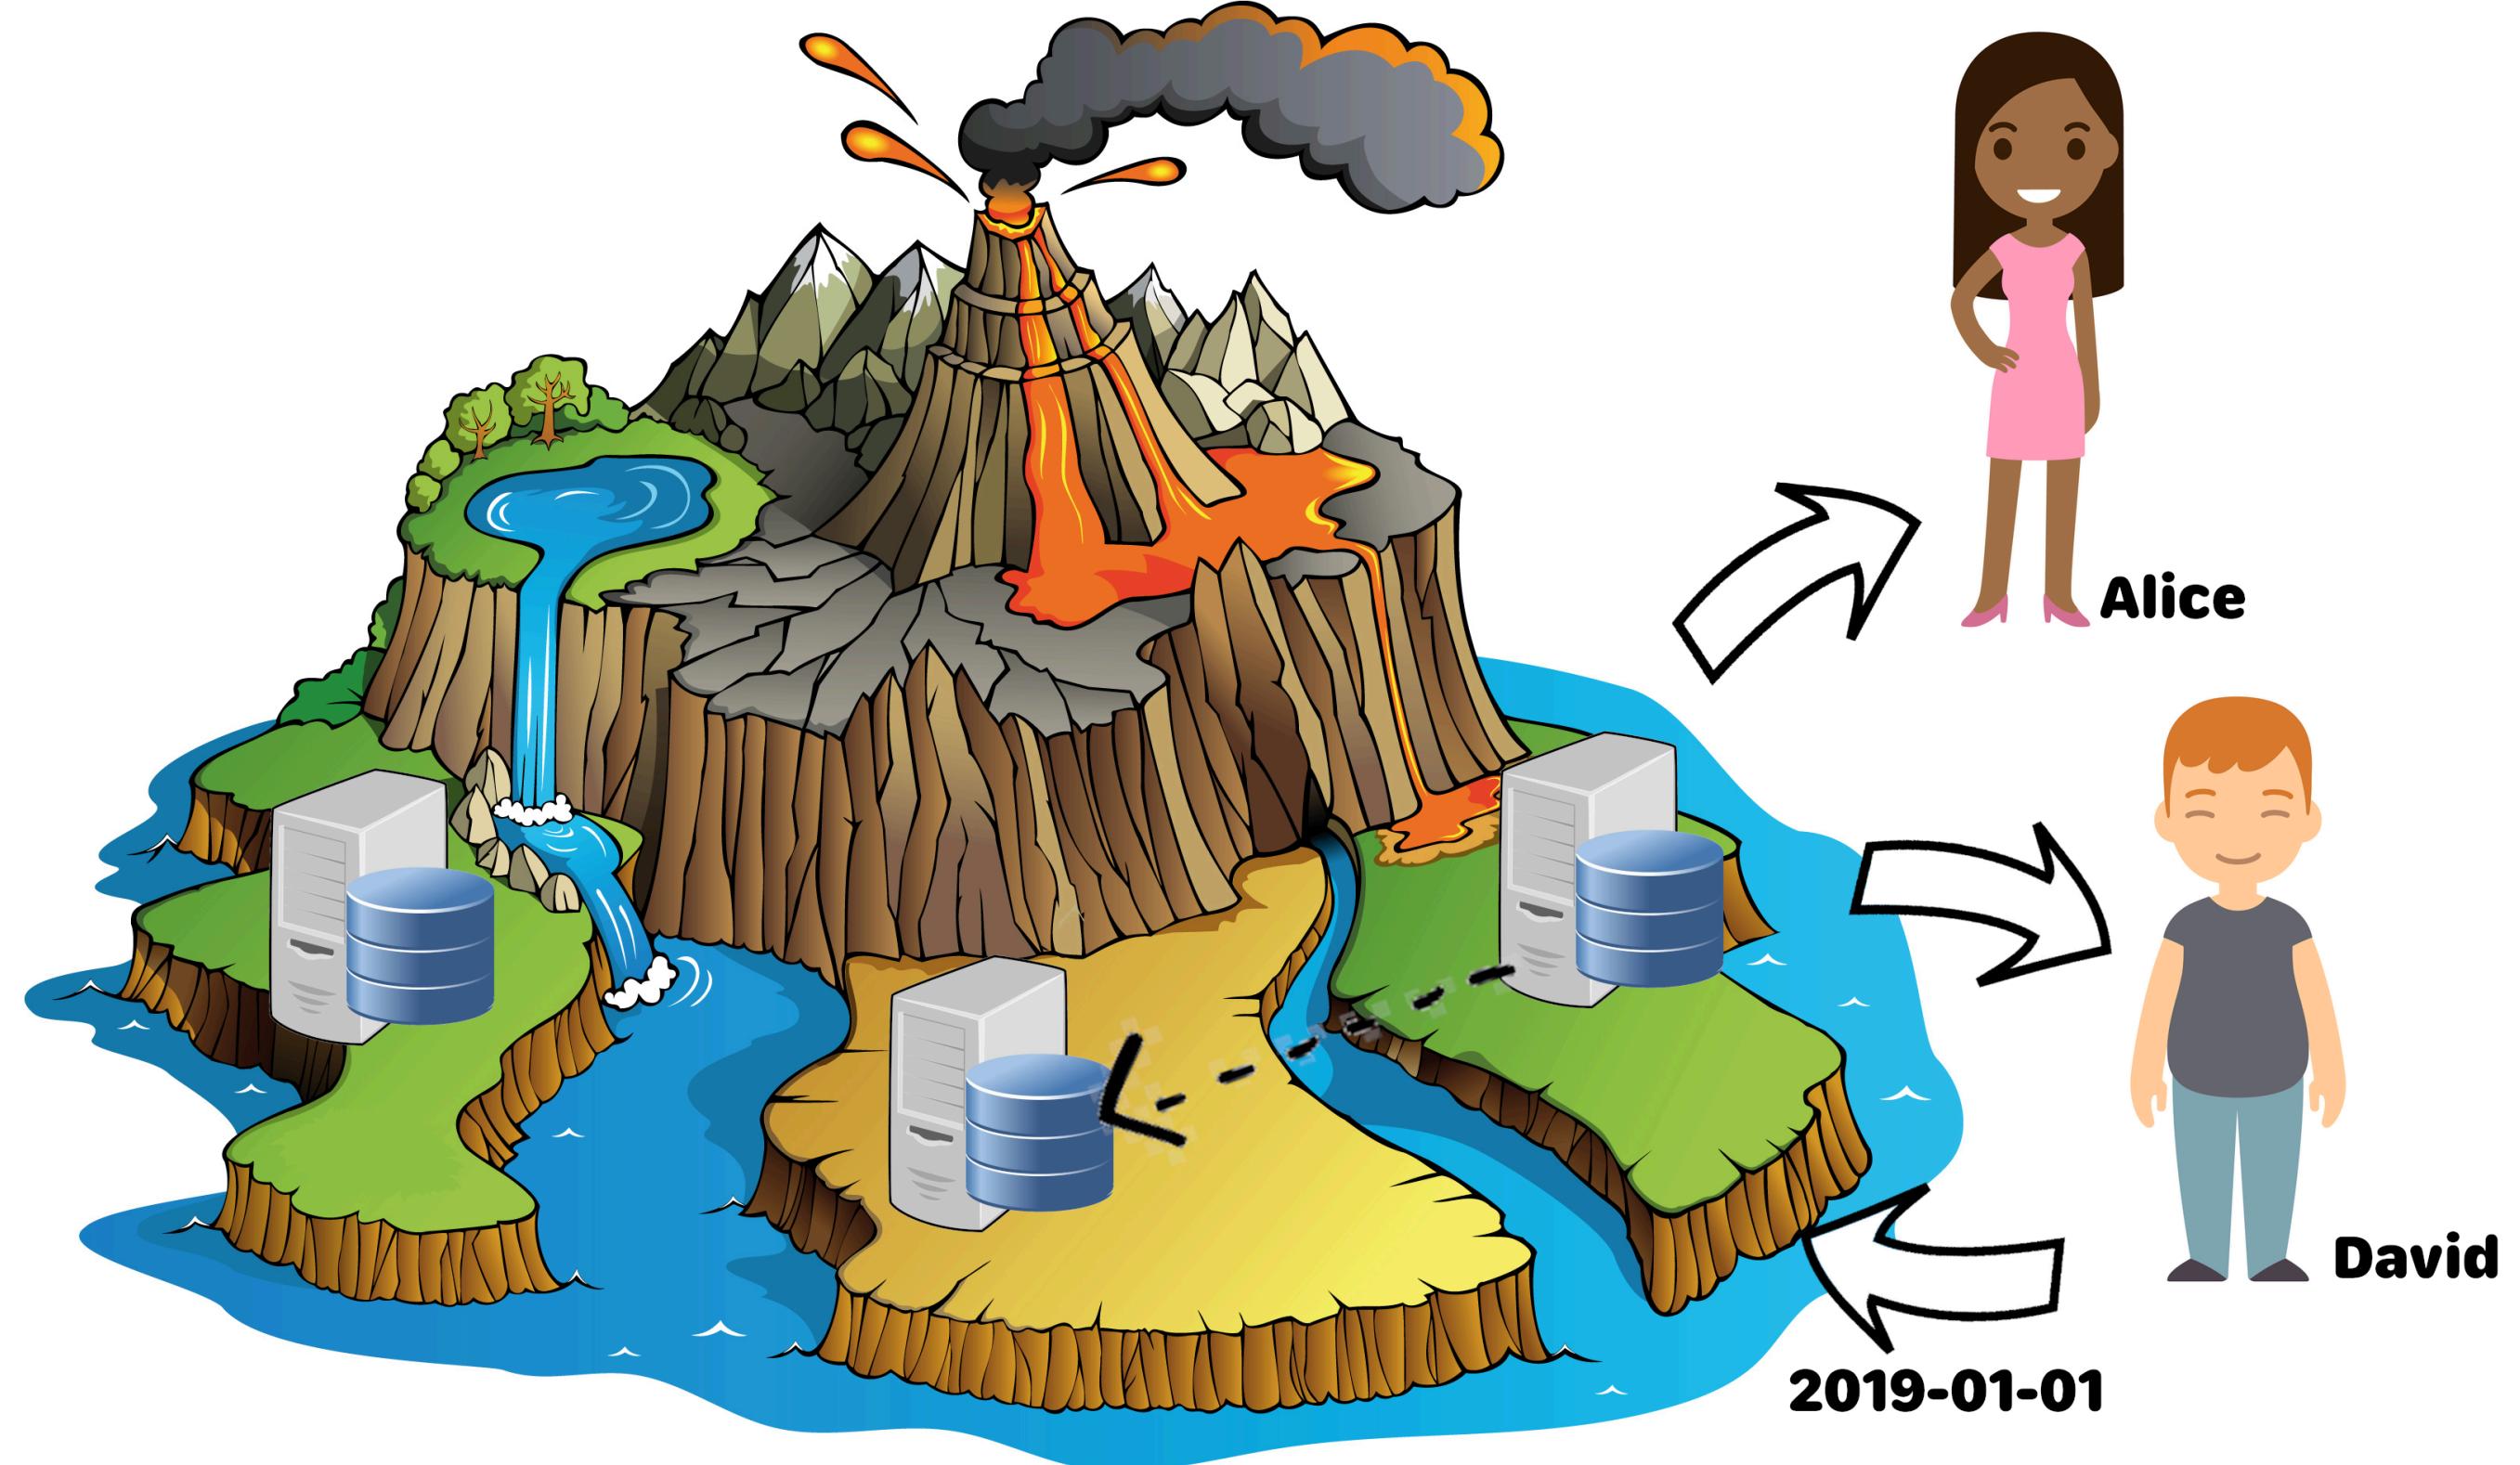
2019-01-04



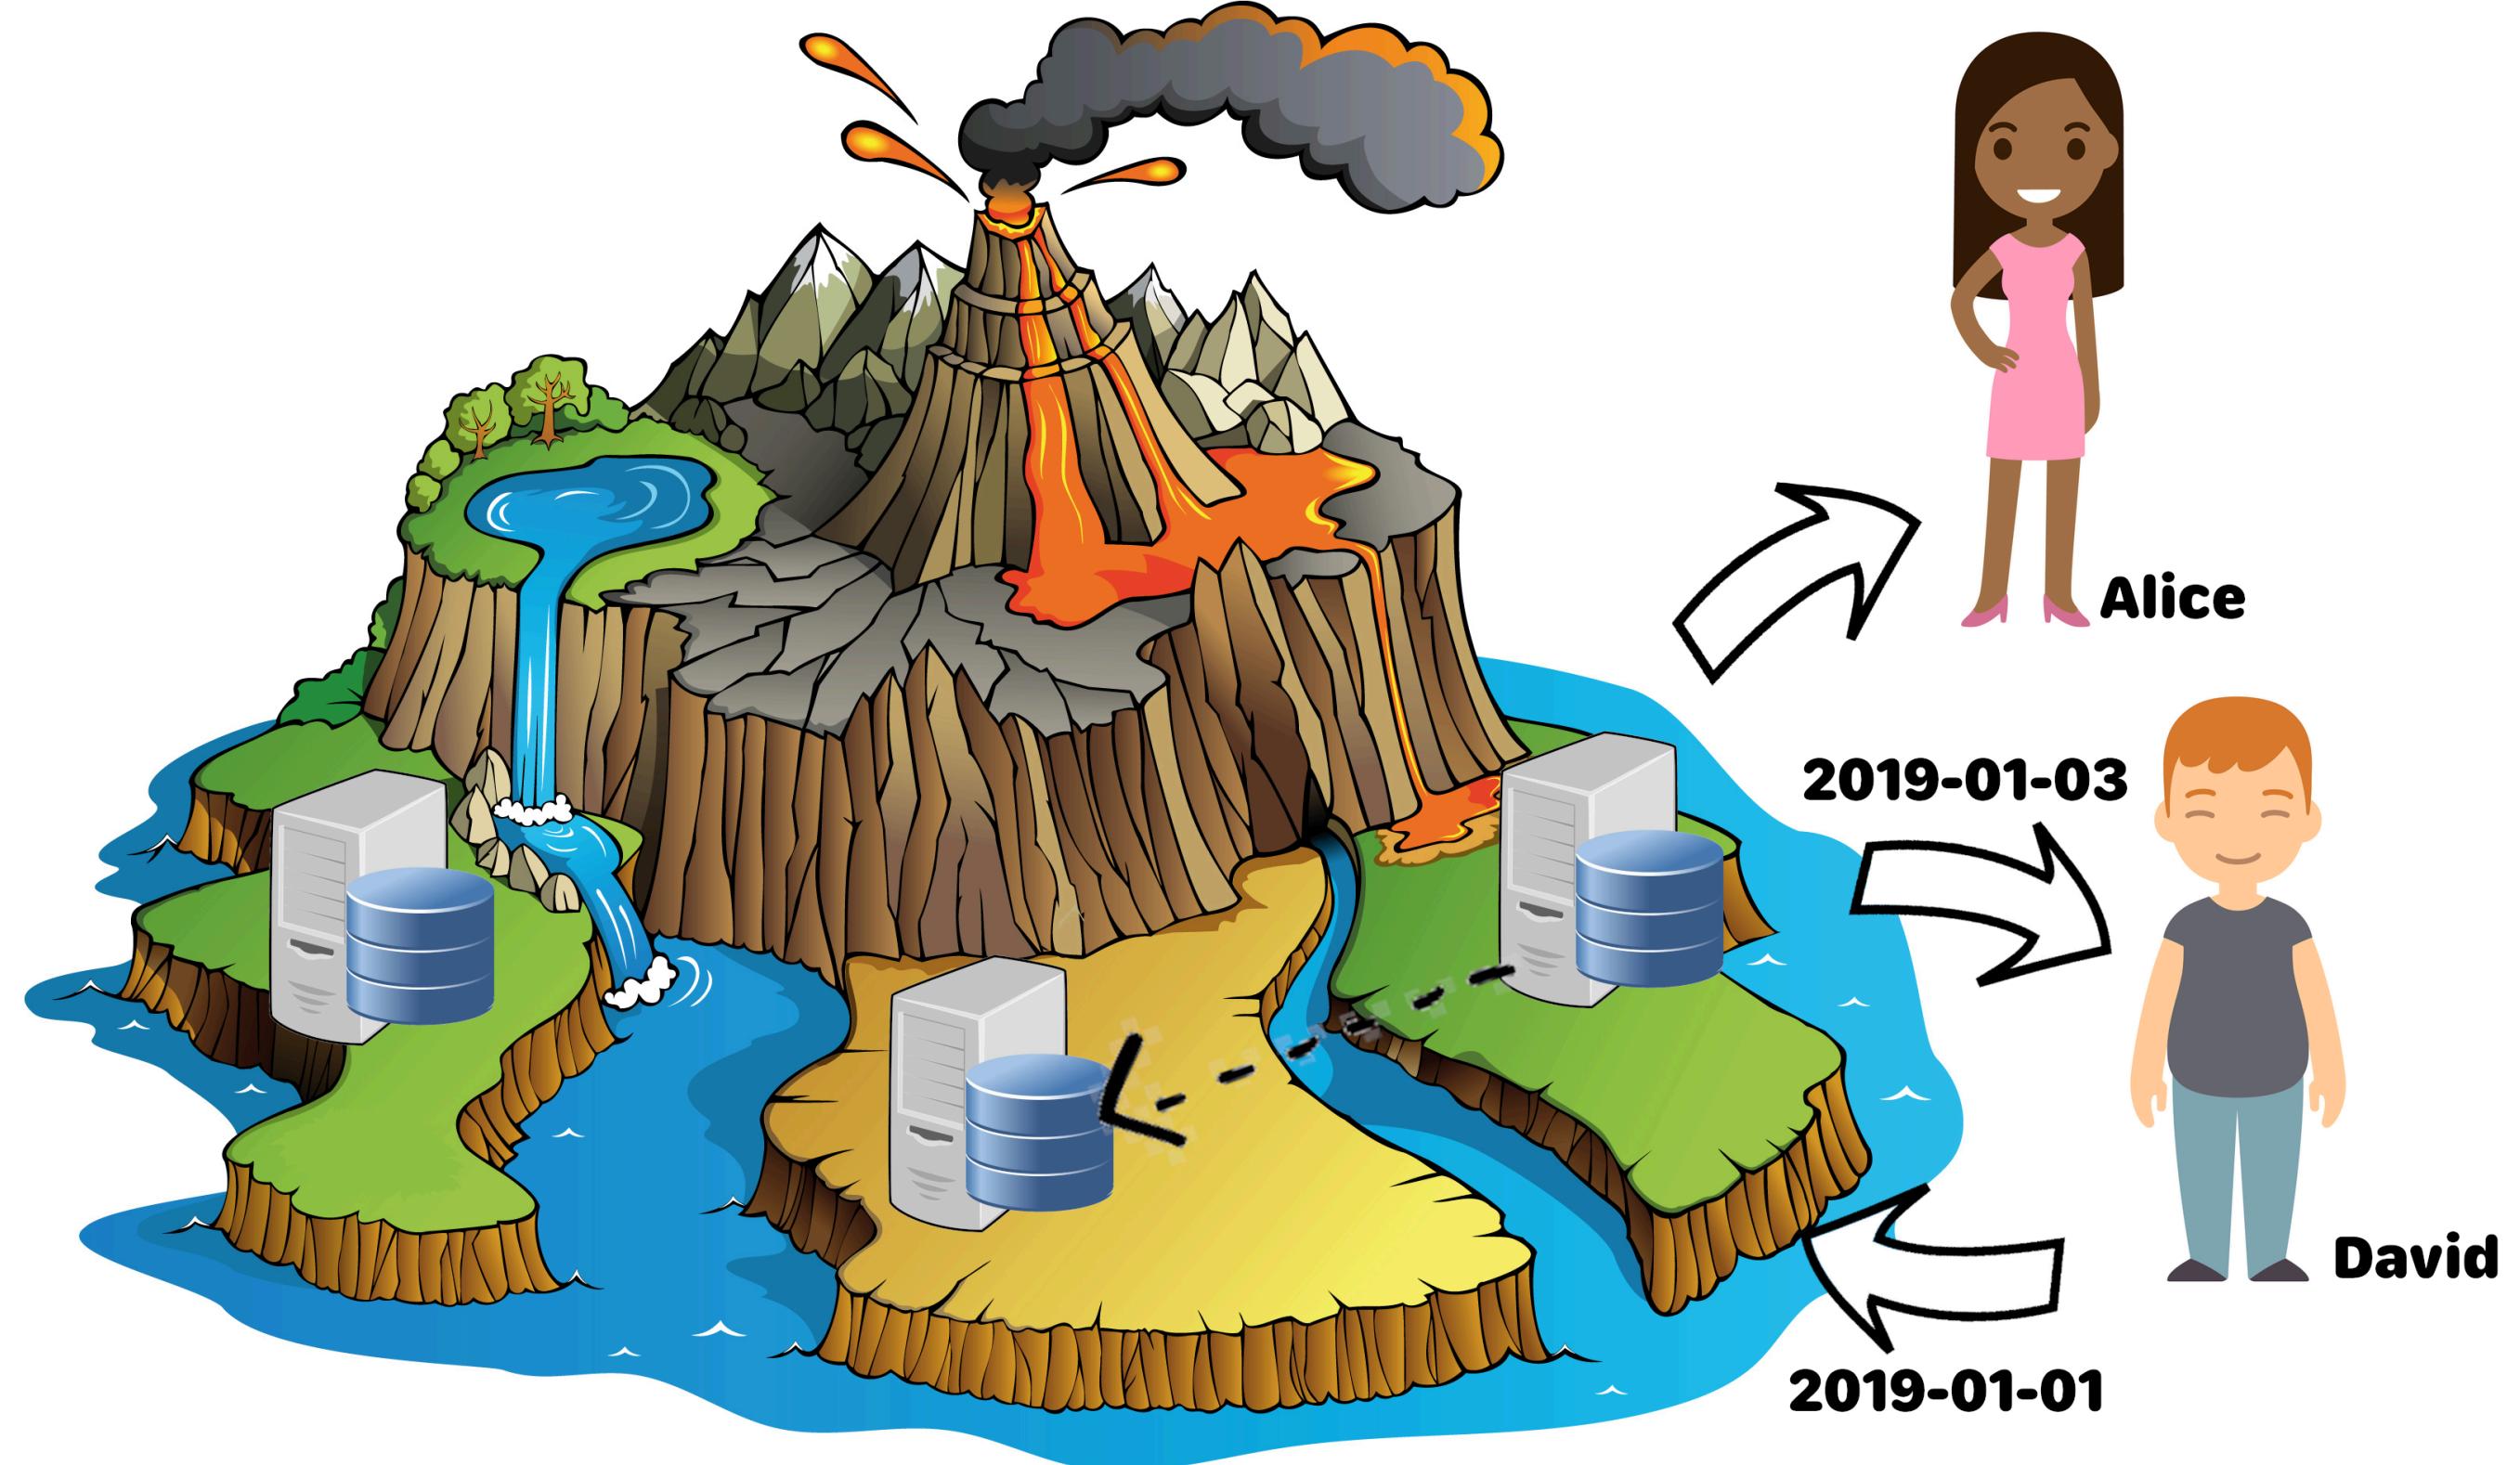
2019-01-04



**2019-01-04**



**2019-01-04**

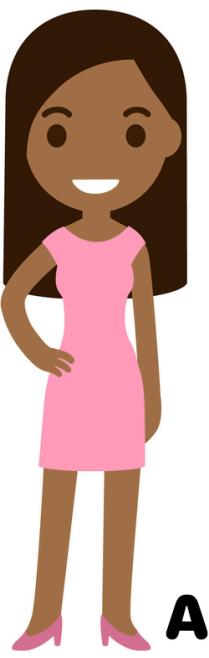


**WHO'S ON THE ISLAND  
ON THE 2ND?**

# **WHO'S ON THE ISLAND?**

- » "as of" - Transaction time
- » "as at" - Valid time

**WHO'S ON THE ISLAND  
ON THE 2ND?  
[AS OF THE 3RD]**



**Alice**

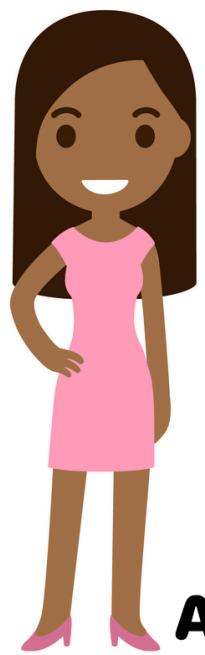


**Bob**



**Carla**

**WHO'S ON THE ISLAND  
ON THE 2ND?  
IAS OF THE 4TH**



**Alice**



**Bob**



**Carla**



**David**

# BITEMPORALITY

## (BOSS SLIDE)

- » Ingest out-of-order temporal data from upstream timestamping systems
- » Reconcile temporal data across eventually consistent systems
- » Dealing with backlog corrections

**AUDIT ALL  
DATA CHANGES AND  
PERFORM DATA FORENSICS**

# **3. STORAGE MODEL**

# **3. STORAGE MODEL**

- » Crux : Document-based
- » Datomic: Entity-Attribute-Value-based

# **CRUX BEING DOCUMENT-BASED MEANS**

# **CRUX BEING DOCUMENT-BASED MEANS**

1. Transactions: apply to the whole doc

# **CRUX BEING DOCUMENT-BASED MEANS**

1. Transactions: apply to the whole doc
2. Whole document gets changed on transactions

# CRUX BEING DOCUMENT-BASED MEANS

1. Transactions: apply to the whole doc
2. Whole document gets changed on transactions
3. Only root-level nodes are indexed

# CRUX BEING DOCUMENT-BASED MEANS

1. Transactions: apply to the whole doc
2. Whole document gets changed on transactions
3. Only root-level nodes are indexed
4. Document relationships can be tricky

# 4. SCHEMA MODEL

# 4. SCHEMA MODEL

- » Crux: No Schema Definition/Enforcement
- » Datomic: Required Schema Definition/Enforcement

# CRUX HAVING NO SCHEMA MEANS

- » No validation (anything goes)
- » Decoupled concerns
- » Silver lining
- » No explicit schema evolution process needed
- » Good practice: implement your own validations

# **5. DEPLOYMENT MODEL**

# 5. DEPLOYMENT MODEL

- » Crux: Self Hosted
  - » Confluent Cloud (KaaS)
  - » Managed hosting available (via Juxt)
- » Datomic:
  - » Datomic On-prem: Self Hosted
  - » Datomic Cloud: Self Hosted (AWS Cloud-based)

# SOME IMPLICATIONS

- » Both can be easily developed locally
  - » run locally, self-host somewhere, or containerize it
- » Crux locally can be a standalone (`crux-core`)
- » Datomic Cloud requires an AWS connection for development

# 6. PRIVACY MODEL

# What is the right to erasure?

Under Article 17 of the GDPR individuals have the right to have personal data erased. This is also known as the 'right to be forgotten'. The right is not absolute and only applies in certain circumstances.

# 6. PRIVACY MODEL

- » Crux: "Eviction" operation
- » Datomic:
  - » Datomic On-prem: "Excision" operation
  - » Datomic Cloud: nothing available

# 7. QUERY DIFFERENCES

# 7. QUERY DIFFERENCES

# 7. QUERY DIFFERENCES

1. Queries are not portable across

# 7. QUERY DIFFERENCES

1. Queries are not portable across
2. Crux does not support Pull Query

# 7. QUERY DIFFERENCES

1. Queries are not portable across
2. Crux does not support Pull Query
3. Datomic does not have Lazy Queries

# 7. QUERY DIFFERENCES

1. Queries are not portable across
2. Crux does not support Pull Query
3. Datomic does not have Lazy Queries
4. Crux does not have sophisticated look up refs

# 7. QUERY DIFFERENCES

1. Queries are not portable across
2. Crux does not support Pull Query
3. Datomic does not have Lazy Queries
4. Crux does not have sophisticated look up refs
5. Datomic requires clause order

# 7. QUERY DIFFERENCES

1. Queries are not portable across
2. Crux does not support Pull Query
3. Datomic does not have Lazy Queries
4. Crux does not have sophisticated look up refs
5. Datomic requires clause order
6. Several other minor differences

# **8. TRANSACTION DIFFERENCES**

# **8. TRANSACTION DIFFERENCES**

# **8. TRANSACTION DIFFERENCES**

1. Transactions are not portable across

# **8. TRANSACTION DIFFERENCES**

1. Transactions are not portable across
2. Crux does not support Components

# **8. TRANSACTION DIFFERENCES**

1. Transactions are not portable across
2. Crux does not support Components
3. Crux does not support Transaction meta model

# **8. TRANSACTION DIFFERENCES**

1. Transactions are not portable across
2. Crux does not support Components
3. Crux does not support Transaction meta model
4. Crux does not support Transaction functions

# **8. TRANSACTION DIFFERENCES**

1. Transactions are not portable across
2. Crux does not support Components
3. Crux does not support Transaction meta model
4. Crux does not support Transaction functions
5. Datomic's injection speed is bound to transactor

# 8. TRANSACTION DIFFERENCES

1. Transactions are not portable across
2. Crux does not support Components
3. Crux does not support Transaction meta model
4. Crux does not support Transaction functions
5. Datomic's injection speed is bound to transactor
6. And remember: doc-based vs. attribute-based

# **9. OTHER ARCHITECTURAL DIFFERENCES**

# 9. OTHER ARCHITECTURAL DIFFERENCES

	<b>CRUX</b>	<b>DATOMIC</b>
Topology	Each Node Storage/ Index	Global Shared Storage/Index
Index Sharding	No Auto Sharding	Peers auto-shard based on working sets
Transactor model	Single Unpartitioned Kafka Topic	Explicit Transactor

# **10. OTHER "FREEBIES"**

# 10. OTHER "FREEBIES"

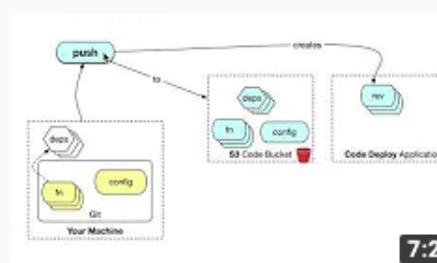
- » Datomic
- » Datomic Cloud
- » Datomic Ions
- » Datomic Cast



## Rich Hickey on Datomic Ions, September 12, 2018

Clojure/nyc • 9.1K views • 11 months ago

The incidental complexity dragon never sleeps! In building Datomic Cloud we took on simplifying Datomic deployment, security ...



## Datomic Ions in Seven Minutes

ClojureTV • 5.5K views • 1 year ago

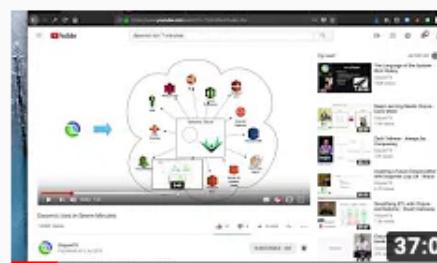
Stuart Halloway introduces Ions for Datomic Cloud on AWS.



## Declarative Domain Modeling for Datomic Ion/Cloud - Tiago Luchini

ClojureTV • 2.6K views • 8 months ago

What if we could build on top of Datomic Ions already easy-to-use, easy-to-deploy, setup and make some common scenarios even ...



## Datomic Ions Hello World in 25 minutes

Tiago Luchini • 1.4K views • 1 year ago

I've set myself the target to see how long it would take to code and deploy a Hello World in Datomic Ions. We cover: - setting up a ...



## Datomic Ions - Stuart Halloway

ClojureTV • 5.8K views • 1 year ago

Datomic Ions (<https://goo.gl/XcEQNh>) let you develop applications for the cloud by deploying your code to a running Datomic ...



## Serverless-ish: Zero to App with Datomic Cloud and GraphQL - Chris Johnson Bidler

ClojureTV • 1.5K views • 8 months ago

The #serverless architecture pattern is taking the world by storm, and for good reasons. The principles of server less design allow ...

# 10. OTHER "FREEBIES"

- » Crux
- » Open Source
- » I.e. Active Object Backend
- » Coming up: Crux console

**CLOJURE IS  
GREATLY SERVED ON  
THE DATABASE SIDE**

# **WITH EITHER CRUX OR DATOMIC:**

1. Immutable Database
2. Query Like a Ninja
3. Unbundle Your Database

# THANKYOU!

@tiagoluchini

<https://luchini.nyc>